

CGI PROGRAMMING 101

Perl for the World Wide Web

2nd Edition

© 2004 by JACQUELINE D. HAMILTON

The following material is excerpted from “CGI Programming 101” (2nd edition) by Jacqueline D. Hamilton. It is copyrighted and may not be redistributed. You may make one printed copy for your own personal use; any other reproduction or retransmission of this document is prohibited.

I ntroduction

This book is intended for web designers, entrepreneurs, students, teachers, and anyone who is interested in learning CGI programming. You do not need any programming experience to get started; if you can write HTML, you can write CGI programs. If you have a website, and want to add guestbook forms, counters, shopping carts, or other interactive elements to your site, then this book is for you.

What is CGI?

“CGI” stands for “Common Gateway Interface.” CGI is one method by which a web server can obtain data from (or send data to) databases, documents, and other programs, and present that data to viewers via the web. More simply, a CGI is a program intended to be run on the web. A CGI program can be written in any programming language, but Perl is one of the most popular, and for this book, Perl is the language we’ll be using.

Why learn CGI?

If you’re going to create web pages, then at some point you’ll want to add a counter, a form to let visitors send you mail or place an order, or something similar. CGI enables you to do that and much more. From mail-forms and counter programs, to the most complex database programs that generate entire websites on-the-fly, CGI programs deliver a broad spectrum of content on the web today.

Why use this book?

This book will get you up and running in as little as a day, teaching you the basics of CGI programs, the fundamentals of Perl, and the basics of processing forms and writing simple programs. Then we’ll move on to advanced topics, such as reading and writing data files, searching for data in files, writing advanced, multi-part forms like order forms and shopping carts, using randomness to spice up your pages, using server-side includes, cookies, and other useful CGI tricks. Things that you’ve probably thought beyond your reach, things you thought you had to pay a programmer to do . . . all of these are things you can easily write yourself, and this book will show you how.

You can also try it out before buying the book; the first six chapters are available online, free of charge, at <http://www.cgi101.com/book/>.

What do you need to get started?

You should already have some experience building web pages and writing HTML. You'll also need Perl and a web server (such as Apache) that is configured to allow you to run your own CGI programs.

The book is written towards CGI programming on Unix, but you can also set up Apache and Perl on Mac OS X and Windows. I've written several online tutorials that will show you how to get started:

Windows: <http://www.cgi101.com/book/connect/windows.html>

How to set up Apache and Perl; how to configure Apache; where to write your programs; differences between CGI programs on Windows and Unix

Mac OS X: <http://www.cgi101.com/book/connect/mac.html>

How to configure Apache (which you already have installed); where to write your programs

Unix: <http://www.cgi101.com/book/connect/unix.html>

How to upload programs to your Unix-based server; Unix tutorial; where to write your programs; Unix permissions.

If you need an ISP that offers CGI hosting, visit <http://www.cgi101.com/hosting>. CGI101 offers Unix shell access, CGI programming, a MySQL database, and all of the Perl modules used in this book. It's an easy, hassle-free way to get started writing your own CGI programs.

Working Code

All of the code examples in this book are available on the web at <http://www.cgi101.com/book/>. You can download any or all of them from there, but do try writing the programs yourself first; you'll learn faster that way.

Conventions Used in this Book

Perl code will be set apart from the text by indenting and use of a fixed-width font:

```
print "This is a print statement.\n";
```

Unix shell commands are shown in a bold font: **chmod 755 filename**

Each program in the book is followed by a link to its source code:

☞ Source code: <http://www.cgi101.com/book/chX/program-cgi.html>

In most cases, a link to a working example is also included:

⇒ Working example: <http://www.cgi101.com/book/chX/demo.html>

Each chapter has its own web page at <http://www.cgi101.com/book/chX>, where X is the chapter number. The full text of chapters 1-6 are online; other chapters include an index of the CGI programs and HTML forms from that chapter, links to online resources mentioned in that chapter, questions and answers relating to the chapter material, plus any chapter errata.

What's New In This Edition?

The 2nd edition of *CGI Programming 101* has been substantially revised from the first edition. You'll learn about Perl modules from the beginning, and work with modules (including the CGI.pm module, which offers many great features for writing CGI programs) throughout the book. You'll learn how to password protect an area on your website, how to build an online catalog with a shopping cart, how to work with cookies, how to protect your site from spammers, and much more.

So turn to Chapter 1, and let's get started.



Getting Started

Our programming language of choice for this book is Perl. Perl is a simple, easy to learn language, yet powerful enough to accomplish very difficult and complex tasks. It is widely available, and is probably already installed on your Unix server. You don't need to compile your Perl programs; you simply write your code, save the file, and run it (or have the web server run it). The program itself is a simple text file; the Perl interpreter does all the work. The advantage to this is you can move your program with little or no changes to any machine with a Perl interpreter. The disadvantage is you won't discover any bugs in your program until you run it.

You can write and edit your CGI programs (which are often called *scripts*) either on your local machine or in the Unix shell. If you're using Unix, try `pico` – it's a very simple, easy to use text editor. Just type **pico filename** to create or edit a file. Type **man pico** for more information and help using `pico`. If you're not familiar with the Unix shell, see Appendix A for a Unix tutorial and command reference.

You can also use a text editor on your local machine and upload the finished programs to the web server. You should either use a plain text editor, such as Notepad (PC) or BBEdit (Mac), or a programming-specific editor that provides some error- and syntax-checking for you. Visit <http://www.cgi101.com/book/editors.html> for a list of some editors you can use to write your CGI programs.

If you use a text editor, be sure to turn off special characters such as “smartquotes.” CGI files must be ordinary text.

Once you've written your program, you'll need to upload it to the web server (unless you're using `pico` and writing it on the server already). You can use any FTP or SCP (secure copy) program to upload your files; a list of some popular FTP and SCP programs can be found at <http://www.cgi101.com/book/connect/>.

It is imperative that you upload your CGI programs as plain text (ASCII) files, and not binary. If you upload your program as a binary file, it may come across with a lot of control characters at the end of the lines, and these will cause errors in your program. You can save yourself a lot of time and grief by just uploading everything as text (unless you're uploading pictures – for example, GIFs or JPEGs – or other true binary data). HTML and Perl CGI programs are not binary, they are plain text.

Once your program is uploaded to the web server, you'll want to be sure to move it to your cgi-bin (or public_html directory – wherever your ISP has told you to put your CGI programs). Then you'll also need to change the permissions on the file so that it is “executable” (or runnable) by the system. The Unix shell command for this is:

chmod 755 filename

This sets the file permissions so that you can read, write, and execute the file, and all other users (including the webserver) can read and execute it. See Appendix A for a full description of **chmod** and its options.

Most FTP and SCP programs allow you to change file permissions; if you use your FTP client to do this, you'll want to be sure that the file is readable and executable by everyone, and writable only by the owner (you).

One final note: Perl code is case-sensitive, as are Unix commands and filenames. Please keep this in mind as you write your first programs, because in Unix “perl” is not the same as “PERL”.

What Is This Unix Shell?

It's a command-line interface to the Unix machine – somewhat like DOS. You have to use a Telnet or SSH (secure shell) program to connect to the shell; see <http://www.cgi101.com/class/connect.html> for a list of some Telnet and SSH programs you can download. Once you're logged in, you can use shell commands to move around, change file permissions, edit files, create directories, move files, and much more.

If you're using a Unix system to learn CGI, you may want to stop here and look at Appendix A to familiarize yourself with the various shell commands. Download a Telnet or SSH program and login to your shell account, then try out some of the commands so you feel comfortable navigating in the shell.

Throughout the rest of this book you'll see Unix shell commands listed in **bold** to set them apart from HTML and CGI code. If you're using a Windows server, you can ignore most of the shell commands, as they don't apply.

Basics of a Perl Program

You should already be familiar with HTML, and so you know that certain things are necessary in the structure of an HTML document, such as the `<head>` and `<body>` tags, and that other tags like links and images have a certain allowed syntax. Perl is very similar; it has a clearly defined syntax, and if you follow those syntax rules, you can write Perl as easily as you do HTML.

The first line of your program should look like this:

```
#!/usr/bin/perl -wT
```

The first part of this line, `#!`, indicates that this is a script. The next part, `/usr/bin/perl`, is the location (or *path*) of the Perl interpreter. If you aren't sure where Perl lives on your system, try typing **which perl** or **whereis perl** in the shell. If the system can find it, it will tell you the full path name to the Perl interpreter. That path is what you should put in the above statement. (If you're using ActivePerl on Windows, the path should be `/perl/bin/perl` instead.)

The final part contains optional flags for the Perl interpreter. Warnings are enabled by the `-w` flag. Special user input taint checking is enabled by the `-T` flag. We'll go into taint checks and program security later, but for now it's good to get in the habit of using both of these flags in all of your programs.

You'll put the text of your program after the above line.

Basics of a CGI Program

A CGI is simply a program that is called by the webserver, in response to some action by a web visitor. This might be something simple like a page counter, or a complex form-handler. Shopping carts and e-commerce sites are driven by CGI programs. So are ad banners; they keep track of who has seen and clicked on an ad.

CGI programs may be written in *any* programming language; we're just using Perl because it's fairly easy to learn. If you're already an expert in some other language and are just reading to get the basics, here it is: if you're writing a CGI that's going to generate an HTML page, you must include this statement somewhere in the program before you print out anything else:

```
print "Content-type: text/html\n\n";
```

This is a content-type header that tells the receiving web browser what sort of data it is about to receive – in this case, an HTML document. If you forget to include it, or if you print something else before printing this header, you’ll get an “Internal Server Error” when you try to access the CGI program.

Your First CGI Program

Now let’s try writing a simple CGI program. Enter the following lines into a new file, and name it “first.cgi”. Note that even though the lines appear indented on this page, you do not have to indent them in your file. The first line (`#!/usr/bin/perl`) should start in column 1. The subsequent lines can start in any column.

Program 1-1: first.cgi

Hello World Program

```
#!/usr/bin/perl -wT
print "Content-type: text/html\n\n";
print "Hello, world!\n";
```

☞ Source code: <http://www.cgi101.com/book/ch1/first.cgi.html>

☞ Working example: <http://www.cgi101.com/book/ch1/first.cgi>

Save (or upload) the file into your web directory, then **chmod 755 first.cgi** to change the file permissions (or use your FTP program to change them). You will have to do this every time you create a *new* program; however, if you’re editing an *existing* program, the permissions will remain the same and shouldn’t need to be changed again.

Now go to your web browser and type the direct URL for your new CGI. For example:

```
http://www.cgi101.com/book/ch1/first.cgi
```

Your actual URL will depend on your ISP. If you have an account on cgi101, your URL is:

```
http://www.cgi101.com/~youruserid/first.cgi
```

You should see a web page with “Hello, world!” on it. (If it you get a “Page Not Found” error, you have the URL wrong. If you got an “Internal Server Error”, see the “Debugging Your Programs,” section at the end of this chapter.)

Let’s try another example. Start a new file (or if you prefer, edit your existing first.cgi) and add some additional print statements. It’s up to your program to print out all of the HTML you want to display in the visitor’s browser, so you’ll have to include print

statements for every HTML tag:

Program 1-2: second.cgi**Hello World Program 2**

```
#!/usr/bin/perl -wT
print "Content-type: text/html\n\n";
print "<html><head><title>Hello World</title></head>\n";
print "<body>\n";
print "<h2>Hello, world!</h2>\n";
print "</body></html>\n";
```

📄 Source code: <http://www.cgi101.com/book/ch1/second-cgi.html>

🔗 Working example: <http://www.cgi101.com/book/ch1/second.cgi>

Save this file, adjust the file permissions if necessary, and view it in your web browser. This time you should see “Hello, world!” displayed in a H2-size HTML header.

Now not only have you learned to write your first CGI program, you’ve also learned your first Perl statement, the `print` function:

```
print "somestring";
```

This function will write out any string, variable, or combinations thereof to the current output channel. In the case of your CGI program, the current output is being printed to the visitor’s browser.

The `\n` you printed at the end of each string is the *newline* character. Newlines are not required, but they will make your program’s output easier to read.

You can write multiple lines of text without using multiple `print` statements by using the here-document syntax:

```
print <<EndMarker;
line1
line2
line3
etc.
EndMarker
```

You can use any word or phrase for the end marker (you’ll see an example next where we use “EndOfHTML” as the marker); just be sure that the closing marker matches the opening marker exactly (it is case-sensitive), and also that the closing marker is on a line by itself, with no spaces before or after the marker.

Let's try it in a CGI program:

Program 1-3: third.cgi**Hello World Program, with here-doc**

```
#!/usr/bin/perl -wT
print "Content-type: text/html\n\n";
print <<EndOfHTML;
<html><head><title>Test Page</title></head>
<body>
<h2>Hello, world!</h2>
</body></html>
EndOfHTML
```

📄 Source code: <http://www.cgi101.com/book/ch1/third-cgi.html>

🔗 Working example: <http://www.cgi101.com/book/ch1/third.cgi>

When a closing here-document marker is on the last line of the file, be sure you have a line break after the marker. If the end-of-file mark is on the same line as the here-doc marker, you'll get an error when you run your program.

The CGI.pm Module

Perl offers a powerful feature to programmers: add-on modules. These are collections of pre-written code that you can use to do all kinds of tasks. You can save yourself the time and trouble of reinventing the wheel by using these modules.

Some modules are included as part of the Perl distribution; these are called *standard library modules* and don't have to be installed. If you have Perl, you already have the standard library modules.

There are also many other modules available that are not part of the standard library. These are typically listed on the Comprehensive Perl Archive Network (CPAN), which you can search on the web at <http://search.cpan.org>.

The CGI.pm module is part of the standard library, and has been since Perl version 5.004. (It should already be installed; if it's not, you either have a very old or very broken version of Perl.) CGI.pm has a number of useful functions and features for writing CGI programs, and its use is preferred by the Perl community. We'll be using it frequently throughout the book.

Let's see how to use a module in your CGI program. First you have to actually include

the module via the `use` command. This goes after the `#!/usr/bin/perl` line and before any other code:

```
use CGI qw(:standard);
```

Note we're not doing `use CGI.pm` but rather `use CGI`. The `.pm` is implied in the `use` statement. The `qw(:standard)` part of this line indicates that we're importing the "standard" set of functions from `CGI.pm`.

Now you can call the various module functions by typing the function name followed by any arguments:

```
functionname(arguments)
```

If you aren't passing any arguments to the function, you can omit the parentheses.

A *function* is a piece of code that performs a specific task; it may also be called a *subroutine* or a *method*. Functions may accept optional *arguments* (also called *parameters*), which are values (strings, numbers, and other variables) passed into the function for it to use. The `CGI.pm` module has many functions; for now we'll start by using these three:

```
header;  
start_html;  
end_html;
```

The `header` function prints out the "Content-type" header. With no arguments, the type is assumed to be "text/html". `start_html` prints out the `<html>`, `<head>`, `<title>` and `<body>` tags. It also accepts optional arguments. If you call `start_html` with only a single string argument, it's assumed to be the page title. For example:

```
print start_html("Hello World");
```

will print out the following*:

```
<html>  
<head>  
<title>Hello World</title>  
<head>  
<body>
```

* Actually `start_html` prints out a full XML header, complete with XML and DOCTYPE tags. In other words, it creates a proper HTML header for your page.

You can also set the page colors and background image with `start_html`:

```
print start_html(-title=>"Hello World",
    -bgcolor=>"#cccccc", -text=>"#999999",
    -background=>"bgimage.jpg");
```

Notice that with multiple arguments, you have to specify the name of each argument with `-title=>`, `-bgcolor=>`, etc. This example generates the same HTML as above, only the body tag indicates the page colors and background image:

```
<body bgcolor="#cccccc" text="#999999"
background="bgimg.jpg">
```

The `end_html` function prints out the closing HTML tags:

```
</body>
</html>
```

The Other Way To Use CGI.pm or “There’s More Than One Way To Do Things In Perl”

As you learn Perl you’ll discover there are often many different ways to accomplish the same task. CGI.pm exemplifies this; it can be used in two different ways. The first way you’ve learned already: function-oriented style. Here you must specify `qw(:standard)` in the use line, but thereafter you can just call the functions directly:

```
use CGI qw(:standard);
print header;
print start_html("Hello World");
```

The other way is *object-oriented* style, where you create an object (or *instance* of the module) and use that to call the various functions of CGI.pm:

```
use CGI;                                # don't need qw(:standard)
$cgi = CGI->new;                        # ($cgi is now the object)
print $cgi->header;                     # function call: $obj->function
print $cgi->start_html("Hello World");
```

Which style you use is up to you. The examples in this book use the function-oriented style, but feel free to use whichever style you’re comfortable with.

So, as you can see, using CGI.pm in your CGI programs will save you some typing. (It also has more important uses, which we'll get into later on.)

Let's try using CGI.pm in an actual program now. Start a new file and enter these lines:

Program 1-4: fourth.cgi**Hello World Program, using CGI.pm**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
print header;
print start_html("Hello World");
print "<h2>Hello, world!</h2>\n";
print end_html;
```

☞ Source code: <http://www.cgi101.com/book/ch1/fourth-cgi.html>

☞ Working example: <http://www.cgi101.com/book/ch1/fourth.cgi>

Be sure to change the file permissions (**chmod 755 fourth.cgi**), then test it out in your browser.

CGI.pm also has a number of functions that serve as HTML shortcuts. For instance:

```
print h2("Hello, world!");
```

Will print an H2-sized header tag. You can find a list of all the CGI.pm functions by typing **perldoc CGI** in the shell, or visiting <http://www.perldoc.com/> and entering “CGI.pm” in the search box.

Documenting Your Programs

Documentation can be embedded in a program using *comments*. A comment in Perl is preceded by the # sign; anything appearing after the # is a comment:

Program 1-5: fifth.cgi**Hello World Program, with Comments**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
# This is a comment
# So is this
#
# Comments are useful for telling the reader
# what's happening. This is important if you
```

```
# write code that someone else will have to
# maintain later.
print header;    # here's a comment. print the header
print start_html("Hello World");
print "<h2>Hello, world!</h2>\n";
print end_html; # print the footer
# the end.
```

☞ Source code: <http://www.cgi101.com/book/ch1/fifth-cgi.html>

⇒ Working example: <http://www.cgi101.com/book/ch1/fifth.cgi>

You'll notice the first line (`#!/usr/bin/perl`) is a comment, but it's a special kind of comment. On Unix, it indicates what program to use to run the rest of the script.

There are several situations in Perl where an `#`-sign is *not* treated as a comment. These depend on specific syntax, and we'll look at them later in the book.

Any line that starts with an `#`-sign is a comment, and you can also put comments at the end of a line of Perl code (as we did in the above example on the `header` and `end_html` lines). Even though comments will only be seen by someone reading the source code of your program, it's a good idea to add comments to your code explaining what's going on. Well-documented programs are much easier to understand and maintain than programs with no documentation.

Debugging Your Programs

A number of problems can happen with your CGI programs, and unfortunately the default response of the webserver when it encounters an error (the "Internal Server Error") is not very useful for figuring out what happened.

If you see the code for the actual Perl program instead of the desired output page from your program, this probably means that your web server isn't properly configured to run CGI programs. You'll need to ask your webmaster how to run CGI programs on your server. And if you ARE the webmaster, check your server's documentation to see how to enable CGI programs.

If you get an Internal Server Error, there's either a permissions problem with the file (did you remember to **chmod 755** the file?) or a bug in your program. A good first step in debugging is to use the `CGI::Carp` module in your program:

```
use CGI::Carp qw(warningsToBrowser fatalToBrowser);
```

This causes all warnings and fatal error messages to be echoed in your browser window. You'll want to remove this line after you're finished developing and debugging your programs, because Carp errors can give away important security info to potential hackers.

If you're using the Carp module and are still seeing the "Internal Server Error", you can further test your program from the command line in the Unix shell. This will check the syntax of your program without actually running it:

```
perl -cwT fourth.cgi
```

If there are errors, it will report any syntax errors in your program:

```
% perl -cwT fourth.cgi  
syntax error at fourth.cgi line 5, near "print"  
fourth.cgi had compilation errors.
```

This tells you there's a problem on or around line 5; make sure you didn't forget a closing semicolon on the previous line, and check for any other typos. Also be sure you saved and uploaded the file as text; hidden control characters or smartquotes can cause syntax errors, too.

Another way to get more info about the error is to look at the webserver log files. Usually this will show you the same information that the CGI::Carp module does, but it's good to know where the server logs are located, and how to look at them. Some usual locations are /usr/local/etc/httpd/logs/error_log, or /var/log/httpd/error_log. Ask your ISP if you aren't sure of the location. In the Unix shell, you can use the **tail** command to view the end of the log file:

```
tail /var/log/apache/error_log
```

The last line of the file should be your error message (although if you're using a shared webserver like an ISP, there will be other users' errors in the file as well). Here are some example errors from the error log:

```
[Fri Jan 16 02:06:10 2004] access to /home/book/ch1/test.cgi failed for  
205.188.198.46, reason: malformed header from script.  
In string, @yahoo now must be written as \@yahoo at /home/book/ch1/test.cgi line  
331, near "@yahoo"  
Execution of /home/book/ch1/test.cgi aborted due to compilation errors.  
[Fri Jan 16 10:04:31 2004] access to /home/book/ch1/test.cgi failed for  
204.87.75.235, reason: Premature end of script headers
```

A "malformed header" or "premature end of script headers" can either mean that you

printed something before printing the “Content-type: text/html” line, or your program died. An error usually appears in the log indicating where the program died, as well.

Resources

The CGI.pm module: <http://stein.cshl.org/WWW/software/CGI/>

The *Official Guide to Programming with CGI.pm*, by Lincoln Stein

Visit <http://www.cgi101.com/book/ch1/> for source code and links from this chapter.

2

Perl Variables

Before you can proceed much further with CGI programming, you'll need some understanding of Perl variables and data types. A *variable* is a place to store a value, so you can refer to it or manipulate it throughout your program. Perl has three types of variables: scalars, arrays, and hashes.

Scalars

A *scalar* variable stores a single (scalar) value. Perl scalar names are prefixed with a dollar sign (\$), so for example, \$x, \$y, \$z, \$username, and \$url are all examples of scalar variable names. Here's how variables are set:

```
$foo = 1;  
$name = "Fred";  
$pi = 3.141592;
```

In this example \$foo, \$name, and \$pi are scalars. You do not have to *declare* a variable before using it, but it's considered good programming style to do so. There are several different ways to declare variables, but the most common way is with the `my` function:

```
my $foo = 1;  
my ($name) = "Fred";  
my ($pi) = 3.141592;
```

`my` simultaneously declares the variables and limits their *scope* (the area of code that can see these variables) to the enclosing code block. (We'll talk more about scope later.) You can declare a variable without giving it a value:

```
my $foo;
```

You can also declare several variables with the same `my` statement:

```
my ($foo, $bar, $blee);
```

You can omit the parentheses if you are declaring a single variable, however a list of variables must be enclosed in parentheses.

A scalar can hold data of any type, be it a string, a number, or whatnot. You can also use scalars in double-quoted strings:

```
my $fnord = 23;
my $blee = "The magic number is $fnord.";
```

Now if you print `$blee`, you will get “The magic number is 23.” Perl *interpolates* the variables in the string, replacing the variable name with the value of that variable.

Let’s try it out in a CGI program. Start a new program called `scalar.cgi`:

Program 2-1: scalar.cgi**Print Scalar Variables Program**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

my $email = "fnord\@cgi101.com";
my $url = "http://www.cgi101.com";

print header;
print start_html("Scalars");
print <<EndHTML;
<h2>Hello</h2>
<p>
My e-mail address is $email, and my web url is
<a href="$url">$url</a>.
</p>
EndHTML

print end_html;
```

📄 Source code: <http://www.cgi101.com/book/ch2/scalar-cgi.html>

🔗 Working example: <http://www.cgi101.com/book/ch2/scalar.cgi>

You may change the `$email` and `$url` variables to show your own e-mail address* and website URL. Save the program, **chmod 755 scalar.cgi**, and test it in your browser.

You'll notice a few new things in this program. First, there's `use strict`. This is a standard Perl module that requires you to declare all variables. You don't *have* to use the `strict` module, but it's considered good programming style, so it's good to get in the habit of using it.

You'll also notice the variable declarations:

```
my $email = "fnord\@cgi101.com";  
my $url = "http://www.cgi101.com";
```

Notice that the `@`-sign in the e-mail address is *escaped* with (preceded by) a backslash. This is because the `@`-sign means something special to Perl – just as the dollar sign indicates a scalar variable, the `@`-sign indicates an array, so if you want to actually use special characters like `@`, `$`, and `%` inside a double-quoted string, you have to precede them with a backslash (`\`).

A better way to do this would be to use a single-quoted string for the e-mail address:

```
my $email = 'fnord@cgi101.com';
```

Single-quoted strings are not interpolated the way double-quoted strings are, so you can freely use the special characters `$`, `@` and `%` in them. However this also means you can't use a single-quoted string to print out a variable, because

```
print '$fnord';
```

will print the actual string `"$fnord"` . . . not the value stored in the variable named `$fnord`.

Arrays

An array stores an ordered list of values. While a scalar variable can only store one value, an array can store many. Perl array names are prefixed with an `@`-sign. Here is an example:

* You should try to avoid leaving your e-mail address permanently displayed on your web site. Spammers routinely crawl the web looking for e-mail addresses. You're better off using a guestbook form. See Chapter 4.

```
my @colors = ("red", "green", "blue");
```

Each individual item (or *element*) of an array may be referred to by its *index* number. Array indices start with 0, so to access the first element of the array `@colors`, you use `$colors[0]`. Notice that when you're referring to a single element of an array, you prefix the name with `$` instead of `@`. The `$`-sign again indicates that it's a *single* (scalar) value; the `@`-sign means you're talking about the *entire* array.

If you want to loop through an array, printing out all of the values, you could print each element one at a time:

```
my @colors = ("red", "green", "blue");

print "$colors[0]\n";      # prints "red"
print "$colors[1]\n";      # prints "green"
print "$colors[2]\n";      # prints "blue"
```

A much easier way to do this is to use a *foreach* loop:

```
my @colors = ("red", "green", "blue");
foreach my $i (@colors) {
    print "$i\n";
}
```

For each iteration of the *foreach* loop, `$i` is set to an element of the `@colors` array. In this example, `$i` is "red" the first time through the loop. The braces `{ }` define where the loop begins and ends, so for any code appearing between the braces, `$i` is set to the current loop iterator.

Notice we've used `my` again here to declare the variables. In the *foreach* loop, `my $i` declares the loop iterator (`$i`) and also limits its scope to the *foreach* loop itself. After the loop completes, `$i` no longer exists.

We'll cover loops more in Chapter 5.

Getting Data Into And Out Of Arrays

An array is an ordered list of elements. You can think of it like a group of people standing in line waiting to buy tickets. Before the line forms, the array is empty:

```
my @people = ();
```

Then Howard walks up. He's the first person in line. To add him to the `@people` array, use the `push` function:

```
push(@people, "Howard");
```

Now Sara, Ken, and Josh get in line. Again they are added to the array using the `push` function. You can push a list of values onto the array:

```
push(@people, ("Sara", "Ken", "Josh"));
```

This pushes the list containing “Sara”, “Ken” and “Josh” onto the end of the `@people` array, so that `@people` now looks like this: (“Howard”, “Sara”, “Ken”, “Josh”)

Now the ticket office opens, and Howard buys his ticket and leaves the line. To remove the first item from the array, use the `shift` function:

```
my $who = shift(@people);
```

This sets `$who` to “Howard”, and also removes “Howard” from the `@people` array, so `@people` now looks like this: (“Sara”, “Ken”, “Josh”)

Suppose Josh gets paged, and has to leave. To remove the last item from the array, use the `pop` function:

```
my $who = pop(@people);
```

This sets `$who` to “Josh”, and `@people` is now (“Sara”, “Ken”)

Both `shift` and `pop` change the array itself, by removing an element from the array.

Finding the Length of Arrays

If you want to find out how many elements are in a given array, you can use the `scalar` function:

```
my @people = ("Howard", "Sara", "Ken", "Josh");
my $linelen = scalar(@people);
print "There are $linelen people in line.\n";
```

This prints “There are 4 people in line.” Of course, there's always more than one way to do things in Perl, and that's true here – the `scalar` function is not actually needed. All you have to do is evaluate the array in a scalar context. You can do this by assigning it to

a scalar variable:

```
my $linelen = @people;
```

This sets `$linelen` to 4.

What if you want to print the name of the last person in line? Remember that Perl array indices start with 0, so the index of the last element in the array is actually `length-1`:

```
print "The last person in line is $people[$linelen-1].\n";
```

Perl also has a handy shortcut for finding the index of the last element of an array, the `$#` shortcut:

```
print "The last person in line is $people[$#people].\n";
```

`$#arrayname` is equivalent to `scalar(@arrayname)-1` . This is often used in `foreach` loops where you loop through an array by its index number:

```
my @colors = ("cyan", "magenta", "yellow", "black");
foreach my $i (0..$#colors) {
    print "color $i is $colors[$i]\n";
}
```

This will print out “color 0 is cyan, color 1 is magenta”, etc.

The `$#arrayname` syntax is one example where an `#` -sign does not indicate a comment.

Array Slices

You can retrieve part of an array by specifying the range of indices to retrieve:

```
my @colors = ("cyan", "magenta", "yellow", "black");
my @slice = $colors[1..2];
```

This example sets `@slice` to (“magenta”, “yellow”).

Finding An Item In An Array

If you want to find out if a particular element exists in an array, you can use the `grep` function:

```
my @results = grep(/pattern/,@listname);
```

`/pattern/` is a *regular expression* for the pattern you’re looking for. It can be a plain string, such as `/Box kite/`, or a complex regular expression pattern.

`/pattern/` will match partial strings inside each array element. To match the entire array element, use `/^pattern$/`, which anchors the pattern match to the beginning (^) and end (\$) of the string. We’ll look more at regular expressions in Chapter 13.

`grep` returns a list of the elements that matched the pattern.

Sorting Arrays

You can do an alphabetical (ASCII) sort on an array of strings using the `sort` function:

```
my @colors = ("cyan", "magenta", "yellow", "black");  
my @colors2 = sort(@colors);
```

`@colors2` becomes the `@colors` array in alphabetically sorted order (“black”, “cyan”, “magenta”, “yellow”). Note that the `sort` function, unlike `push` and `pop`, does *not* change the original array. If you want to save the sorted array, you have to assign it to a variable. If you want to save it back to the original array variable, you’d do:

```
@colors = sort @colors;
```

You can invert the order of the array with the `reverse` function:

```
my @colors = ("cyan", "magenta", "yellow", "black");  
@colors = reverse(@colors);
```

`@colors` is now (“black”, “yellow”, “magenta”, “cyan”).

To do a reverse sort, use both functions:

```
my @colors = ("cyan", "magenta", "yellow", "black");  
@colors = reverse(sort(@colors));
```

`@colors` is now (“yellow”, “magenta”, “cyan”, “black”).

The `sort` function, by default, compares the ASCII values of the array elements (see <http://www.cgi101.com/book/ch2/ascii.html> for the chart of ASCII values). This means if you try to sort a list of numbers, you get “12” before “2”. You can do a true numeric sort like so:

```
my @numberlist = (8, 4, 3, 12, 7, 15, 5);
my @sortednumberlist = sort( {$a <=> $b;} @numberlist);
```

`{ $a <=> $b; }` is actually a small subroutine, embedded right in your code, that gets called for each pair of items in the array. It compares the first number (`$a`) to the second number (`$b`) and returns a number indicating whether `$a` is greater than, equal to, or less than `$b`. This is done repeatedly with all the numbers in the array until the array is completely sorted.

We'll talk more about custom sorting subroutines in Chapter 12.

Joining Array Elements Into A String

You can merge an array into a single string using the `join` function:

```
my @colors = ("cyan", "magenta", "yellow", "black");
my $colorstring = join(", ", @colors);
```

This joins `@colors` into a single string variable (`$colorstring`), with each element of the `@colors` array combined and separated by a comma and a space. In this example `$colorstring` becomes “cyan, magenta, yellow, black”.

You can use any string (including the empty string) as the separator. The separator is the first argument to the `join` function:

```
join(separator, list);
```

The opposite of `join` is `split`, which splits a string into a list of values. See Chapter 7 for more on `split`.

Array or List?

In general, any function or syntax that works for arrays will also work for a list of values:

```
my $color = ("red", "green", "blue")[1];
# $color is "green"

my $colorstring = join(", ", ("red", "green", "blue"));
# $colorstring is now "red, green, blue"

my ($first, $second, $third) = sort("red", "green", "blue");
# $first is "blue", $second is "green", $third is "red"
```


Hashes

A *hash* is a special kind of array – an *associative* array, or paired list of elements. Each pair consists of a string *key* and a *data value*.

Perl hash names are prefixed with a percent sign (%). Here's how they're defined:

Hash Name	key	value
my %colors = ("red",	"#ff0000",
	"green",	"#00ff00",
	"blue",	"#0000ff",
	"black",	"#000000",
	"white",	"#ffffff");

This particular example creates a hash named %colors which stores the RGB HEX values for the named colors. The color names are the hash keys; the hex codes are the hash values.

Remember that there's more than one way to do things in Perl, and here's the other way to define the same hash:

```
my %colors = ( red    => "#ff0000",
               green  => "#00ff00",
               blue   => "#0000ff",
               black  => "#000000",
               white  => "#ffffff" );
```

The => operator automatically quotes the left side of the argument, so enclosing quotes around the key names are not needed.

To refer to the individual elements of the hash, you'll do:

```
$colors{'red'}
```

Here, "red" is the key, and \$colors{'red'} is the value associated with that key. In this case, the value is "#ff0000".

You don't usually need the enclosing quotes around the value, either; \$colors{red} also works if the key name doesn't contain characters that are also Perl operators (things like +, -, =, * and /).

To print out all the values in a hash, you can use a foreach loop:

```
foreach my $color (keys %colors) {
    print "$colors{$color}=$color\n";
}
```

This example uses the `keys` function, which returns a list of the keys of the named hash. One drawback is that `keys %hashname` will return the keys in unpredictable order – in this example, `keys %colors` could return (“red”, “blue”, “green”, “black”, “white”) or (“red”, “white”, “green”, “black”, “blue”) or any combination thereof. If you want to print out the hash in exact order, you have to specify the keys in the `foreach` loop:

```
foreach my $color ("red", "green", "blue", "black", "white") {
    print "$colors{$color}=$color\n";
}
```

Let’s write a CGI program using the colors hash. Start a new file called `colors.cgi`:

Program 2-2: colors.cgi	Print Hash Variables Program
--------------------------------	-------------------------------------

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

# declare the colors hash:
my %colors = ( red => "#ff0000", green=> "#00ff00",
               blue => "#0000ff", black => "#000000",
               white => "#ffffff" );

# print the html headers
print header;
print start_html("Colors");

foreach my $color (keys %colors) {
    print "<font color=\"${colors{$color}}\">$color</font>\n";
}
print end_html;
```

📄 Source code: <http://www.cgi101.com/book/ch2/colors-cgi.html>

🔗 Working example: <http://www.cgi101.com/book/ch2/colors.cgi>

Save it and **chmod 755 colors.cgi**, then test it in your web browser.

Notice we’ve had to add backslashes to escape the quotes in this double-quoted string:

```
print "<font color=\"${colors}${color}\">${color}</font>\n";
```

A better way to do this is to use Perl's qq operator:

```
print qq(<font color="${colors}${colors}">${color}</font>\n);
```

qq creates a double-quoted string for you. And it's much easier to read without all those backslashes in there.

Adding Items to a Hash

To add a new value to a hash, you simply do:

```
$hashname{newkey} = newvalue;
```

Using our colors example again, here's how to add a new value with the key "purple":

```
$colors{purple} = "#ff00ff";
```

If the named key already exists in the hash, then an assignment like this overwrites the previous value associated with that key.

Determining Whether an Item Exists in a Hash

You can use the `exists` function to see if a particular key/value pair exists in the hash:

```
exists $hashname{key}
```

This returns a true or false value. Here's an example of it in use:

```
if (exists $colors{purple}) {  
    print "Sorry, the color purple is already in the  
hash.<br>\n";  
} else {  
    $colors{purple} = "#ff00ff";  
}
```

This checks to see if the key "purple" is already in the hash; if not, it adds it.

Deleting Items From a Hash

You can delete an individual key/value pair from a hash with the `delete` function:

```
delete $hashname{key};
```

If you want to empty out the entire hash, do:

```
%hashname = ();
```

Values

We've already seen that the `keys` function returns a list of the keys of a given hash. Similarly, the `values` function returns a list of the hash values:

```
my %colors = (red => "#ff0000", green=> "#00ff00",
              blue => "#0000ff", black => "#000000",
              white => "#ffffff" );

my @keyslice = keys %colors;
# @keyslice now equals a randomly ordered list of
# the hash keys:
# ("red", "green", "blue", "black", "white")

my @valueslice = values %colors;
# @valueslice now equals a randomly ordered list of
# the hash values:
# ("ff0000", "#00ff00", "#0000ff", "#000000", "#ffffff")
```

As with `keys`, `values` returns the values in unpredictable order.

Determining Whether a Hash is Empty

You can use the `scalar` function on hashes as well:

```
scalar($hashname);
```

This returns true or false value – true if the hash contains any key/value pairs. The value returned does *not* indicate how many pairs are in the hash, however. If you want to find that number, use:

```
scalar keys(%hashname);
```

Here's an example:

```
my %colors = (red => "#ff0000", green=> "#00ff00",
              blue => "#0000ff", black => "#000000",
              white => "#ffffff" );

my $numcolors = scalar(keys(%colors));
print "There are $numcolors in this hash.\n";
```

This will print out “There are 5 colors in this hash.”

Resources

Visit <http://www.cgi101.com/book/ch2/> for source code and links from this chapter.

Chapter 2 Reference: Arrays and Hashes

Array Functions

<code>@array = ();</code>	Defines an empty array
<code>@array = ("a", "b", "c");</code>	Defines an array with values
<code>\$array[0]</code>	The first element of <code>@array</code>
<code>\$array[0] = a;</code>	Sets the first element of <code>@array</code> to <i>a</i>
<code>\$array[3..5]</code>	Array slice - returns a list containing the 3rd thru 5th elements of <code>@array</code>
<code>scalar(@array)</code>	Returns the number of elements in <code>@array</code>
<code>\$#array</code>	The index of the last element in <code>@array</code>
<code>grep(/pattern/, @array)</code>	Returns a list of the items in <code>@array</code> that matched <code>/pattern/</code>
<code>join(expr, @array)</code>	Joins <code>@array</code> into a single string separated by <i>expr</i>
<code>push(@array, \$var)</code>	Adds <code>\$var</code> to <code>@array</code>
<code>pop(@array)</code>	Removes last element of <code>@array</code> and returns it
<code>reverse(@array)</code>	Returns <code>@array</code> in reverse order
<code>shift(@array)</code>	Removes first element of <code>@array</code> and returns it
<code>sort(@array)</code>	Returns alphabetically sorted <code>@array</code>
<code>sort({\$a<=>\$b}, @array)</code>	Returns numerically sorted <code>@array</code>

Hash Functions

<code>%hash = ();</code>	Defines an empty hash
<code>%hash = (a => 1, b => 2);</code>	Defines a hash with values
<code>\$hash{\$key}</code>	The value referred to by this <code>\$key</code>
<code>\$hash{\$key} = \$value;</code>	Sets the value referred to by <code>\$key</code>
<code>exists \$hash{\$key}</code>	True if the key/value pair exists
<code>delete \$hash{\$key}</code>	Deletes the key/value pair specified by <code>\$key</code>
<code>keys %hash</code>	Returns a list of the hash keys
<code>values %hash</code>	Returns a list of the hash values



CGI Environment Variables

Environment variables are a series of hidden values that the web server sends to every CGI program you run. Your program can parse them and use the data they send. Environment variables are stored in a hash named %ENV:

Key	Value
DOCUMENT_ROOT	The root directory of your server
HTTP_COOKIE	The visitor's cookie, if one is set
HTTP_HOST	The hostname of the page being attempted
HTTP_REFERER	The URL of the page that called your program
HTTP_USER_AGENT	The browser type of the visitor
HTTPS	"on" if the program is being called through a secure server
PATH	The system path your server is running under
QUERY_STRING	The query string (see GET, below)
REMOTE_ADDR	The IP address of the visitor
REMOTE_HOST	The hostname of the visitor (if your server has reverse-name-lookups on; otherwise this is the IP address again)
REMOTE_PORT	The port the visitor is connected to on the web server
REMOTE_USER	The visitor's username (for .htaccess-protected pages)
REQUEST_METHOD	GET or POST
REQUEST_URI	The interpreted pathname of the requested document or CGI (relative to the document root)
SCRIPT_FILENAME	The full pathname of the current CGI
SCRIPT_NAME	The interpreted pathname of the current CGI (relative to the document root)
SERVER_ADMIN	The email address for your server's webmaster
SERVER_NAME	Your server's fully qualified domain name (e.g. www.cgi101.com)
SERVER_PORT	The port number your server is listening on
SERVER_SOFTWARE	The server software you're using (e.g. Apache 1.3)

Some servers set other environment variables as well; check your server documentation for more information. Notice that some environment variables give information about your server, and will never change (such as `SERVER_NAME` and `SERVER_ADMIN`), while others give information about the visitor, and will be different every time someone accesses the program.

Not all environment variables get set. `REMOTE_USER` is only set for pages in a directory or subdirectory that's password-protected via a `.htaccess` file. (See Chapter 20 to learn how to password protect a directory.) And even then, `REMOTE_USER` will be the username as it appears in the `.htaccess` file; it's not the person's email address. There is no reliable way to get a person's email address, short of asking them for it with a web form.

You can print the environment variables the same way you would any hash value:

```
print "Caller = $ENV{HTTP_REFERER}\n";
```

Let's try printing some environment variables. Start a new file named `env.cgi`:

Program 3-1: env.cgi**Print Environment Variables Program**

```
#!/usr/bin/perl -wT
use strict;
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);

print header;
print start_html("Environment");

foreach my $key (sort(keys(%ENV))) {
    print "$key = $ENV{$key}<br>\n";
}

print end_html;
```

☞ Source code: <http://www.cgi101.com/book/ch3/env-cgi.html>

⇒ Working example: <http://www.cgi101.com/book/ch3/env.cgi>

Save the file, **chmod 755 env.cgi**, then try it in your web browser. Compare the environment variables displayed with the list on the previous page. Notice which values show information about your server and CGI program, and which ones give away information about you (such as your browser type, computer operating system, and IP address).

Let's look at several ways to use some of this data.

Referring Page

When you click on a hyperlink on a web page, you're being referred to another page. The web server for the receiving page keeps track of the referring page, and you can access the URL for that page via the HTTP_REFERER environment variable. Here's an example:

Program 3-2: refer.cgi**HTTP Referrer Program**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

print header;
print start_html("Referring Page");
print "Welcome, I see you've just come from
$ENV{HTTP_REFERER}!<p>\n";

print end_html;
```

☞ Source code: <http://www.cgi101.com/book/ch3/refer-cgi.html>

⇒ Working example: <http://www.cgi101.com/book/ch3/> (click on refer.cgi)

Remember, HTTP_REFERER only gets set when a visitor actually clicks on a link to your page. If they type the URL directly (or use a bookmarked URL), then HTTP_REFERER is blank. To properly test your program, create an HTML page with a link to refer.cgi, then click on the link:

```
<a href="refer.cgi">Referring Page</a>
```

HTTP_REFERER is not a foolproof method of determining what page is accessing your program. It can easily be forged.

Remote Host Name, and Hostname Lookups

You've probably seen web pages that greet you with a message like "Hello, visitor from (yourhost)!", where (yourhost) is the hostname or IP address you're currently logged in with. This is a pretty easy thing to do because your IP address is stored in the %ENV hash.

If your web server is configured to do hostname lookups, then you can access the visitor's actual hostname from the `$ENV{REMOTE_HOST}` value. Servers often don't do hostname lookups automatically, though, because it slows down the server. Since `$ENV{REMOTE_ADDR}` contains the visitor's IP address, you can reverse-lookup the hostname from the IP address using the `Socket` module in Perl. As with `CGI.pm`, you have to use the `Socket` module:

```
use Socket;
```

(There is no need to add `qw(:standard)` for the `Socket` module.)

The `Socket` module offers numerous functions for socket programming (most of which are beyond the scope of this book). We're only interested in the reverse-IP lookup for now, though. Here's how to do the reverse lookup:

```
my $ip = "209.189.198.102";
my $hostname = gethostbyaddr(inet_aton($ip), AF_INET);
```

There are actually two functions being called here: `gethostbyaddr` and `inet_aton`. `gethostbyaddr` is a built-in Perl function that returns the hostname for a particular IP address. However, it requires the IP address be passed to it in a packed 4-byte format. The `Socket` module's `inet_aton` function does this for you.

Let's try it in a CGI program. Start a new file called `rhost.cgi`, and enter the following code:

Program 3-3: <code>rhost.cgi</code>	Remote Host Program
--	----------------------------

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;
use Socket;

print header;
print start_html("Remote Host");

my $hostname = gethostbyaddr(inet_aton($ENV{REMOTE_ADDR}),
AF_INET);
print "Welcome, visitor from $hostname!<p>\n";

print end_html;
```

☞ Source code: <http://www.cgi101.com/book/ch3/rhost-cgi.html>

⇒ Working example: <http://www.cgi101.com/book/ch3/rhost.cgi>

Detecting Browser Type

The HTTP_USER_AGENT environment variable contains a string identifying the browser (or “user agent”) accessing the page. Unfortunately there is no standard (yet) for user agent strings, so you will see a vast assortment of different strings. Here’s a sampling of some:

```
DoCoMo/1.0/P502i/c10 (Google CHTML Proxy/1.0)
Firefly/1.0 (compatible; Mozilla 4.0; MSIE 5.5)
Googlebot/2.1 (+http://www.googlebot.com/bot.html)
Mozilla/3.0 (compatible)
Mozilla/4.0 (compatible; MSIE 4.01; MSIECrawler; Windows 95)
Mozilla/4.0 (compatible; MSIE 5.0; MSN 2.5; AOL 8.0; Windows 98; DigExt)
Mozilla/4.0 (compatible; MSIE 5.0; Mac_PowerPC)
Mozilla/4.0 (compatible; MSIE 5.0; Windows 98; DigExt; Hotbar 4.1.7.0)
Mozilla/4.0 (compatible; MSIE 6.0; AOL 9.0; Windows NT 5.1)
Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.0; DigExt)
Mozilla/4.0 WebTV/2.6 (compatible; MSIE 4.0)
Mozilla/5.0 (Macintosh; U; PPC Mac OS X; en-US; rv:1.0.2) Gecko/20020924 AOL/7.0
Mozilla/5.0 (Macintosh; U; PPC Mac OS X; en-US; rv:1.0.2) Gecko/20021120 Netscape/
7.01
Mozilla/5.0 (Macintosh; U; PPC Mac OS X; en-us) AppleWebKit/85 (KHTML, like Gecko)
Safari/85
Mozilla/5.0 (Windows; U; Win98; en-US; m18) Gecko/20010131 Netscape6/6.01
Mozilla/5.0 (Slurp/cat; slurp@inktomi.com; http://www.inktomi.com/slurp.html)
Mozilla/5.0 (X11; U; Linux i686; en-US; rv:1.5a) Gecko/20030718
Mozilla/5.0 (compatible; Konqueror/3.0-rc3; i686 Linux; 20020913)
NetNewsWire/1.0 (Mac OS X; Pro; http://ranchero.com/netnewswire/)
Opera/6.0 (Windows 98; U) [en]
Opera/7.10 (Linux 2.4.19 i686; U) [en]
Scooter/3.3
```

As you can see, sometimes the user agent string reveals what type of browser and computer the visitor is using, and sometimes it doesn’t. Some of these aren’t even browsers at all, like the search engine robots (Googlebot, Inktomi and Scooter) and RSS reader (NetNewsWire). You should be careful about writing programs (and websites) that do browser detection. It’s one thing to collect browser info for logging purposes; it’s quite another to design your entire site exclusively for a certain browser. Visitors will be annoyed if they can’t access your site because you think they have the “wrong” browser.

That said, here's an example of how to detect the browser type. This program uses Perl's `index` function to see if a particular substring (such as "MSIE") exists in the `HTTP_USER_AGENT` string. `index` is used like so:

```
index(string, substring);
```

It returns a numeric value indicating where in the string the substring appears, or -1 if the substring does not appear in the string. We use an if/else block in this program to see if the index is greater than -1.

Program 3-4: browser.cgi**Browser Detection Program**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

print header;
print start_html("Browser Detect");

my($ua) = $ENV{HTTP_USER_AGENT};

print "User-agent: $ua<p>\n";
if (index($ua, "MSIE") > -1) {
    print "Your browser is Internet Explorer.<p>\n";
} elsif (index($ua, "Netscape") > -1) {
    print "Your browser is Netscape.<p>\n";
} elsif (index($ua, "Safari") > -1) {
    print "Your browser is Safari.<p>\n";
} elsif (index($ua, "Opera") > -1) {
    print "Your browser is Opera.<p>\n";
} elsif (index($ua, "Mozilla") > -1) {
    print "Your browser is probably Mozilla.<p>\n";
} else {
    print "I give up, I can't tell what browser you're
using!<p>\n";
}

print end_html;
```

📄 Source code: <http://www.cgi101.com/book/ch3/browser-cgi.html>

🔗 Working example: <http://www.cgi101.com/book/ch3/browser.cgi>

If you have several different browsers installed on your computer, try testing the program

with each of them.

We'll look more at if/else blocks in Chapter 5.

A Simple Form Using GET

There are two ways to send data from a web form to a CGI program: GET and POST. These *methods* determine how the form data is sent to the server.

With the GET method, the input values from the form are sent as part of the URL and saved in the QUERY_STRING environment variable. With the POST method, data is sent as an input stream to the program. We'll cover POST in the next chapter, but for now, let's look at GET.

You can set the QUERY_STRING value in a number of ways. For example, here are a number of direct links to the env.cgi program:

```
http://www.cgi101.com/book/ch3/env.cgi?test1
http://www.cgi101.com/book/ch3/env.cgi?test2
http://www.cgi101.com/book/ch3/env.cgi?test3
```

Try opening each of these in your web browser. Notice that the value for QUERY_STRING is set to whatever appears after the question mark in the URL itself. In the above examples, it's set to "test1", "test2", and "test3" respectively.

You can also process simple forms using the GET method. Start a new HTML document called envform.html, and enter this form:

Program 3-5: envform.html	Simple HTML Form Using GET
----------------------------------	-----------------------------------

```
<html><head><title>Test Form</title></head>
<body>

<form action="env.cgi" method="GET">
Enter some text here:
<input type="text" name="sample_text" size=30>
<input type="submit"><p>
</form>

</body></html>
```

⇒ Working example: <http://www.cgi101.com/book/ch3/envform.html>

Save the form and upload it to your website. Remember you may need to change the path to `env.cgi` depending on your server; if your CGI programs live in a “`cgi-bin`” directory then you should use `action="cgi-bin/env.cgi"`.

Bring up the form in your browser, then type something into the input field and hit return. You’ll notice that the value for `QUERY_STRING` now looks like this:

```
sample_text=whatever+you+typed
```

The string to the left of the equals sign is the name of the form field. The string to the right is whatever you typed into the input box. Notice that any spaces in the string you typed have been replaced with a `+`. Similarly, various punctuation and other special non-alphanumeric characters have been replaced with a `%`-code. This is called URL-encoding, and it happens with data submitted through either GET or POST methods.

You can send multiple input data values with GET:

```
<form action="env.cgi" method="GET">
First Name: <input type="text" name="fname" size=30><p>
Last Name: <input type="text" name="lname" size=30><p>
<input type="submit">
</form>
```

This will be passed to the `env.cgi` program as follows:

```
$ENV{QUERY_STRING} = "fname=joe&lname=smith"
```

The two form values are separated by an ampersand (`&`). You can divide the query string with Perl’s `split` function:

```
my @values = split(/&/,$ENV{QUERY_STRING});
```

`split` lets you break up a string into a list of strings, splitting on a specific character. In this case, we’ve split on the “`&`” character. This gives us an array named `@values` containing two elements: (`"fname=joe"`, `"lname=smith"`). We can further split each string on the “`=`” character using a `foreach` loop:

```
foreach my $i (@values) {
    my($fieldname, $data) = split(/=/, $i);
    print "$fieldname = $data<br>\n";
}
```

This prints out the field names and the data entered into each field in the form. It does not

do URL-decoding, however. A better way to parse QUERY_STRING variables is with CGI.pm.

Using CGI.pm to Parse the Query String

If you're sending more than one value in the query string, it's best to use CGI.pm to parse it. This requires that your query string be of the form:

```
fieldname1=value1
```

For multiple values, it should look like this:

```
fieldname1=value1&fieldname2=value2&fieldname3=value3
```

This will be the case if you are using a form, but if you're typing the URL directly then you need to be sure to use a fieldname, an equals sign, then the field value.

CGI.pm provides these values to you automatically with the `param` function:

```
param('fieldname');
```

This returns the value entered in the `fieldname` field. It also does the URL-decoding for you, so you get the exact string that was typed in the form field.

You can get a list of all the fieldnames used in the form by calling `param` with no arguments:

```
my @fieldnames = param();
```

param is NOT a Variable

`param` is a function call. You can't do this:

```
print "$p = param($p)<br>\n";
```

If you want to print the value of `param($p)`, you can print it by itself:

```
print param($p);
```

Or call `param` outside of the double-quoted strings:

```
print "$p = ", param($p), "<br>\n";
```

You won't be able to use `param('fieldname')` inside a here-document. You may find it easier to assign the form values to individual variables:

```
my $firstname = param('firstname');
my $lastname = param('lastname');
```

Another way would be to assign every form value to a hash:

```
my(%form);
foreach my $p (param()) {
    $form{$p} = param($p);
}
```

You can achieve the same result by using CGI.pm's `Vars` function:

```
use CGI qw(:standard Vars);
my %form = Vars();
```

The `Vars` function is not part of the “standard” set of CGI.pm functions, so it must be included specifically in the `use` statement.

Either way, after storing the field values in the `%form` hash, you can refer to the individual field names by using `$form{'fieldname'}`. (This will **not** work if you have a form with multiple fields having the same field name.)

Let's try it now. Create a new form called `getform.html`:

Program 3-6: getform.html	Another HTML Form Using GET
----------------------------------	------------------------------------

```
<html><head><title>Test Form</title></head>
<body>

<form action="get.cgi" method="GET">
First Name: <input type="text" name="firstname" size=30><br>
Last Name: <input type="text" name="lastname" size=30><br>
<input type="submit"><p>
</form>

</body></html>
```

⇒ Working example: <http://www.cgi101.com/book/ch3/getform.html>

Save and upload it to your webserver, then bring up the form in your web browser.

Now create the CGI program called `get.cgi`:

Program 3-7: <code>get.cgi</code>	Form Processing Program Using GET
--	--

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

print header;
print start_html("Get Form");

my %form;
foreach my $p (param()) {
    $form{$p} = param($p);
    print "$p = $form{$p}<br>\n";
}
print end_html;
```

📄 Source code: <http://www.cgi101.com/book/ch3/get-cgi.html>

Save and **`chmod 755 get.cgi`**. Now fill out the form in your browser and press submit. If you encounter errors, refer back to Chapter 1 for debugging.

Take a look at the full URL of `get.cgi` after you press submit. You should see all of your form field names and the data you typed in as part of the URL. This is one reason why GET is not the best method for handling forms; it isn't secure.

GET is NOT Secure

GET is not a secure method of sending data. Don't use it for forms that send password info, credit card data or other sensitive information. Since the data is passed through as part of the URL, it'll show up in the web server's logfile (complete with all the data). Server logfiles are often readable by other users on the system. URL history is also saved in the browser and can be viewed by anyone with access to the computer. Private information should always be sent with the POST method, which we'll cover in the next chapter. (And if you're asking visitors to send sensitive information like credit card numbers, you should also be using a secure server in addition to the POST method.)

There may also be limits to how much data can be sent with GET. While the HTTP protocol doesn't specify a limit to the length of a URL, certain web browsers and/or

servers may.

Despite this, the GET method is often the best choice for certain types of applications. For example, if you have a database of articles, each with a unique article ID, you would probably want a single `article.cgi` program to serve up the articles. With the article ID passed in by the GET method, the program would simply look at the query string to figure out which article to display:

```
<a href="article.cgi?id=22">Article Name</a>
```

We'll be revisiting that idea later in the book. For now, let's move on to Chapter 4 where we'll see how to process forms using the POST method.

Resources

Visit <http://www.cgi101.com/book/ch3/> for source code and links from this chapter.

4

Processing Forms and Sending Mail

Most forms you create will send their data using the POST method. POST is more secure than GET, since the data isn't sent as part of the URL, and you can send more data with POST. Also, your browser, web server, or proxy server may cache GET queries, but posted data is resent each time.

Your web browser, when sending form data, encodes the data being sent. Alphanumeric characters are sent as themselves; spaces are converted to plus signs (+); other characters – like tabs, quotes, etc. – are converted to “%HH” – a percent sign and two hexadecimal digits representing the ASCII code of the character. This is called URL encoding.

In order to do anything useful with the data, your program must decode these. Fortunately the CGI.pm module does this work for you. You access the decoded form values the same way you did with GET:

```
$value = param('fieldname');
```

So you already know how to process forms! You can try it now by changing your getform.html form to method=“POST” (rather than method=“GET”). You'll see that it works identically whether you use GET or POST. Even though the data is sent differently, CGI.pm handles it for you automatically.

The Old Way of Decoding Form Data

Before CGI.pm was bundled with Perl, CGI programmers had to write their own form-parsing code. If you read some older CGI books (including the first edition of this book), or if you're debugging old code, you'll probably encounter the old way of decoding form data. Here's what it looks like:

```

read(STDIN, $buffer, $ENV{'CONTENT_LENGTH'});
@pairs = split(/&/, $buffer);
foreach $pair (@pairs) {
    ($name, $value) = split(/=/, $pair);
    $value =~ tr/+// /;
    $value =~ s/%([a-fA-F0-9][a-fA-F0-9])/pack("C",
hex($1))/eg;
    $FORM{$name} = $value;
}

```

This code block reads the posted form data from standard input, loops through the fieldname=value fields in the form, and uses the pack function to do URL-decoding. Then it stores each fieldname/value pair in a hash called %FORM.

This code is deprecated and should be avoided; use CGI.pm instead. If you want to upgrade an old program that uses the above code block, you can replace it with this:

```

my %FORM;
foreach my $field (param()) {
    $FORM{$field} = param($field);
}

```

Or you could use the Vars function:

```

use CGI qw(:standard Vars);
my %FORM = Vars();

```

Either method will replace the old form-parsing code, although keep in mind that this will not work if your form has multiple fields with the same name. We'll look at how to handle those in the next chapter.

Guestbook Form

One of the first CGI programs you're likely to want to add to your website is a guestbook program, so let's start writing one. First create your HTML form. The actual fields can be up to you, but a bare minimum might look like this:

```

<form action="post.cgi" method="POST">
Your Name: <input type="text" name="name"><br>
Email Address: <input type="text" name="email"><br>
Comments:<br>
<textarea name="comments" rows="5"
    cols="60"></textarea><br>
<input type="submit" value="Send">

```

</form>

☞ Source code: <http://www.cgi101.com/book/ch4/guestbook1.html>

(Stylistically it's better NOT to include a “reset” button on forms like this. It's unlikely the visitor will want to erase what they've typed, and more likely they'll accidentally hit “reset” instead of “send”, which can be an aggravating experience. They may not bother to re-fill the form in such cases.)

Now you need to create `post.cgi`. This is nearly identical to the `get.cgi` from last chapter, so you may just want to copy that program and make changes:

Program 4-1: <code>post.cgi</code>	Form Processing Program Using POST
---	---

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

print header;
print start_html("Thank You");
print h2("Thank You");

my %form;
foreach my $p (param()) {
    $form{$p} = param($p);
    print "$p = $form{$p}<br>\n";
}
print end_html;
```

☞ Source code: <http://www.cgi101.com/book/ch4/post-cgi.html>

⇒ Working example: <http://www.cgi101.com/book/ch4/form.html>

Test your program by entering some data into the fields, and pressing “send” when finished. Notice that the data is not sent in the URL this time, as it was with the GET example.

Of course, this form doesn't actually DO anything with the data, which doesn't make it much of a guestbook. Let's see how to send the data in e-mail.

Sending Mail

There are several ways to send mail. We'll be using the **sendmail** program for these

examples. If you're using a non-Unix system (or a Unix without sendmail installed), there are a number of third-party Perl modules that you can use to achieve the same effect. See <http://search.cpan.org/> (search for "sendmail") for a list of platform-independent mailers, and Chapter 14 for examples of how to install third-party modules. If you're using ActivePerl on Windows, visit <http://www.cgi101.com/book/ch4/> for a link to more information about sending mail from Windows.

Before you can write your form-to-mail CGI program, you'll need to figure out where the sendmail program is installed on your webserver. (For cgi101.com, it's in `/usr/sbin/sendmail`. If you're not sure where it is, try doing **which sendmail** or **whereis sendmail**; usually one of these two commands will yield the correct location.)

Since we're using the `-T` flag for taint checking, the first thing you need to do before connecting to sendmail is set the `PATH` environment variable:

```
$ENV{PATH} = "/usr/sbin";
```

The path should be the directory where sendmail is located; if sendmail is in `/usr/sbin/sendmail`, then `$ENV{PATH}` should be `"/usr/sbin"`. If it's in `/var/lib/sendmail`, then `$ENV{PATH}` should be `"/var/lib"`.

Next you open a *pipe* to the sendmail program:

```
open (MAIL, "|/usr/sbin/sendmail -t -oi") or  
die "Can't fork for sendmail: $!\n";
```

The pipe (which is indicated by the `|` character) causes all of the output printed to that filehandle (`MAIL`) to be fed directly to the `/usr/sbin/sendmail` program as if it were standard input to that program. Several flags are also passed to sendmail:

```
-t      Read message for recipients. To:, Cc:, and Bcc:  
        lines will be scanned for recipient addresses  
-oi     Ignore dots alone on lines by themselves in  
        incoming messages.
```

The `-t` flag tells sendmail to look at the message headers to determine who the mail is being sent to. You'll have to print all of the message headers yourself:

```
my $recipient = 'recipient@cgi101.com';  
  
print MAIL "From: sender@cgi101.com\n";  
print MAIL "To: $recipient\n";  
print MAIL "Subject: Guestbook Form\n\n";
```

Remember that you can safely put an @-sign inside a single-quoted string, like 'recipient@cgi101.com', or you can escape the @-sign in double-quoted strings by using a backslash ("sender\\@cgi101.com").

The message headers are complete when you print a single blank line following the header lines. We've accomplished this by printing two newlines at the end of the subject header:

```
print MAIL "Subject: Guestbook Form\\n\\n";
```

After that, you can print the body of your message.

Let's try it. Start a new file named guestbook.cgi, and edit it as follows. You don't need to include the comments in the following code; they are just there to show you what's happening.

Program 4-2: guestbook.cgi**Guestbook Program**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

print header;
print start_html("Results");

# Set the PATH environment variable to the same path
# where sendmail is located:

$ENV{PATH} = "/usr/sbin";

# open the pipe to sendmail
open (MAIL, "|/usr/sbin/sendmail -oi -t") or
    &dienice("Can't fork for sendmail: $!\n");

# change this to your own e-mail address
my $recipient = 'recipient@cgi101.com';

# Start printing the mail headers
# You must specify who it's to, or it won't be delivered:

print MAIL "To: $recipient\\n";
```

```

# From should probably be the webserver.

print MAIL "From: nobody\@cgi101.com\n";

# print a subject line so you know it's from your form cgi.

print MAIL "Subject: Form Data\n\n";

# Now print the body of your mail message.
foreach my $p (param()) {
    print MAIL "$p = ", param($p), "\n";
}

# Be sure to close the MAIL input stream so that the
# message actually gets mailed.

close(MAIL);

# Now print a thank-you page

print <<EndHTML;
<h2>Thank You</h2>
<p>Thank you for writing!</p>
<p>Return to our <a href="index.html">home page</a>.</p>
EndHTML

print end_html;

# The dienice subroutine handles errors.

sub dienice {
    my($errmsg) = @_;
    print "<h2>Error</h2>\n";
    print "<p>$errmsg</p>\n";
    print end_html;
    exit;
}

```

☞ Source code: <http://www.cgi101.com/book/ch4/guestbook-cgi.html>

Save and chmod the file, then modify your guestbook.html form so that the action points to guestbook.cgi:

```
<form action="guestbook.cgi" method="POST">
```

Try testing the form. If the program runs successfully, you'll get e-mail in a few moments

with the results of your post. (Remember to change `$recipient` to your email address!)

Subroutines

In the guestbook program we used a new structure: a subroutine called “dienice.” A *subroutine* is a user-defined function. You’ve already used functions like `param` and `start_html` from the `CGI.pm` module, and built-in functions like `shift` and `pop`. You can also define your own custom functions.

In the mail program, the `dienice` subroutine is only called if the program can’t open the pipe to sendmail. Rather than aborting and giving you a server error (or worse, NO error), you want your program to give you some useful data about what went wrong; `dienice` does that, by printing the error message and closing HTML tags, and exiting the program. We’ll be using the `dienice` subroutine throughout the rest of the book, as a generic catch-all error-handler.

Subroutines are useful for isolating blocks of code that are reused frequently in your program. The structure of a subroutine is as follows:

```
sub subname {  
    # your code here  
}
```

The subroutine block starts with the word `sub`, followed by the name of the subroutine. The code for the subroutine is then enclosed in curly braces `{ }`.

Subroutines can be placed anywhere in your program, though for readability it’s usually best to put them at the end, after the main program code.

To invoke a subroutine, enter the subroutine name and an optional list of arguments:

```
subname;  
subname(arguments);
```

You may prefix the subroutine name with an `&`-sign:

```
&subname;  
&subname(arguments);
```

The `&`-sign is optional. However, we’ll be using this syntax throughout the book to differentiate calls to subroutines we’ve written ourselves. Calls to built-in functions or functions provided by external modules will not have this sign.

Here is an example of a call to a subroutine named “mysub” with three arguments:

```
&mysub($arg1, "whatever", 23);
```

The arguments are passed to the subroutine in the special Perl array `@_`. You can then assign the elements of that array to special temporary variables, like so:

```
sub mysub {  
    my($arg1, $arg2, $arg3) = @_;  
    # your code here  
}
```

In this example, the `my` function limits the scope of `$arg1`, `$arg2` and `$arg3` to the `mysub` subroutine. This keeps your temporary variables visible only to the subroutine itself (where they’re actually needed and used), rather than to the entire program (where they’re not needed). This also means if you change one of the variables inside your subroutine, the value of the original variable won’t change (unless it’s a reference, which we’ll look at next).

Passing Arrays and Hashes to Subroutines

When passing an array (or a hash) to a subroutine, the array is expanded into a list of its values. This might be okay if the array is the only argument:

```
&subname(@array1);
```

However if you have multiple arguments, you’re going to run into problems:

```
&subname(@array1, $item2, $item3);  
  
sub subname {  
    my(@ary, $arg2, $arg3) = @_;  
}
```

In this example, *all* of the arguments (including `$item2` and `$item3`) are stored in `@ary`, and `$arg2` and `$arg3` are undefined. In order to pass the array or hash properly to the subroutine, you need to pass it as a *reference*, by prefixing the `@` (or `%`) by a backslash:

```
&subname(\@array1, $item2, \%hash1);  
  
sub subname {
```

```
    my($arrayref, $arg2, $hashref) = @_;  
}
```

Now `$arrayref` is a reference to `@array1`, `$arg2` is whatever the value of `$item2` is, and `$hashref` is a reference to `%hash1`. To access individual elements of an array reference, instead of using `$arrayref[1]`, you use `$arrayref->[1]`. Similarly with a hash reference you use `$hashref->{key}` instead of `$hashref{key}`.

A *reference* is a pointer to the original variable. If you change the value of an element of an array reference, you're changing the original array's values.

Optionally you could *dereference* the array inside your subroutine by doing:

```
my @localary = @{$arrayref};
```

A hash is dereferenced like so:

```
my %localhash = %{$hashref};
```

A dereferenced array (or hash) is localized to your subroutine, so you can change the values of `@newarray` or `%newhash` without altering the original variables.

You can find out a lot more about references by reading **perldoc perlref** and **perldoc perlreftut** (the Perl reference tutorial).

Subroutine Return Values

Subroutines can return a value:

```
sub subname {  
    # your code here  
    return $somevalue;  
}
```

If you omit the return statement, then the value returned by the subroutine is the value of the last expression executed in that routine.

If you want to save the return value, be sure to assign it to a variable:

```
my $result = &subname(arguments);
```

Subroutines can also return a list:

```
sub subname {  
    # your code here  
    return $value1, $value2, 'foo';  
}
```

Which can then be assigned to a list of variables:

```
my ($x, $y, $z) = &subname(arguments);
```

Or an array:

```
my @x = &subname(arguments);
```

Return vs. Exit

You'll notice that our `dienice` subroutine does not return a value at all, but rather calls the `exit` function. `exit` causes the entire program to terminate immediately.

Sendmail Subroutine

Here is an example of the mail-sending code in a compact subroutine:

```
sub sendmail {  
    my ($from, $to, $subject, $message) = @_;  
    $ENV{PATH} = "/usr/sbin";  
    open (MAIL, "|/usr/sbin/sendmail -oi -t") or  
        &dienice("Can't fork for sendmail: $!\n");  
    print MAIL "To: $to\n";  
    print MAIL "From: $from\n";  
    print MAIL "Subject: $subject\n\n";  
    print MAIL "$message\n";  
    close(MAIL);  
}
```

Sending Mail to More Than One Recipient

If you want to send mail to more than one email address, just add the desired addresses to the `$recipient` line, separated by commas:

```
my $recipient = 'recipient1@cgi101.com,  
recipient2@cgi101.com, recipient3@cgi101.com';
```

Defending Against Spammers

When building form-to-mail programs, you need to take precautions to prevent spammers from hijacking your programs to send unwanted e-mail to other recipients. They can do this by writing their own form (or program) to send data to *your* CGI. If your program prints any of the form fields as mail headers without checking them first, the spammer can insert their *own* mail headers (and even their own message). The end result: your program becomes a relay for spammers.

The primary defense against this is to not allow the form to specify ANY of the mail headers (such as the From, To, or Subject headers). Note that in our guestbook program, the From, To and Subject headers were all hardcoded in the program.

Of course, it would be nice to have the “From” header show the poster’s e-mail address. You could allow this if you validate it first, verifying that it’s really an e-mail address and doesn’t contain any extra headers. You can validate e-mail addresses by using a regular expression pattern match, which we’ll cover in Chapter 13, or by using the Email::Valid module, which we’ll look at in Chapter 14.

Resources

Visit <http://www.cgi101.com/book/ch4/> for source code and links from this chapter.



Advanced Forms and Perl Control Structures

In the last chapter you learned how to decode form data, and mail it to yourself. However, one problem with the guestbook program is that it didn't do any error-checking or specialized processing. You might not want to get blank forms, or you may want to require certain fields to be filled out. You might also want to write a quiz or questionnaire, and have your program take different actions depending on the answers. All of these things require some more advanced processing of the form data, and that will usually involve using *control structures* in your Perl code.

Control structures include conditional statements, such as `if/elsif/else` blocks, as well as loops like `foreach`, `for` and `while`.

If Conditions

You've already seen `if/elsif` in action. The structure is always started by the word `if`, followed by a condition to be evaluated, then a pair of braces indicating the beginning and end of the code to be executed if the condition is true. The condition is enclosed in parentheses:

```
if (condition) {  
    code to be executed  
}
```

The condition statement can be anything that evaluates to true or false. In Perl, any string is true except the empty string and 0. Any number is true except 0. An undefined value (or `undef`) is false. You can also test whether a certain value equals something, or doesn't equal something, or is greater than or less than something. There are different conditional test operators, depending on whether the variable you want to test is a string or a number:

Relational and Equality Operators

Test	Numbers	Strings
\$x is equal to \$y	<code>\$x == \$y</code>	<code>\$x eq \$y</code>
\$x is not equal to \$y	<code>\$x != \$y</code>	<code>\$x ne \$y</code>
\$x is greater than \$y	<code>\$x > \$y</code>	<code>\$x gt \$y</code>
\$x is greater than or equal to \$y	<code>\$x >= \$y</code>	<code>\$x ge \$y</code>
\$x is less than \$y	<code>\$x < \$y</code>	<code>\$x lt \$y</code>
\$x is less than or equal to \$y	<code>\$x <= \$y</code>	<code>\$x le \$y</code>

If it's a string test, you use the letter operators (`eq`, `ne`, `lt`, etc.), and if it's a numeric test, you use the symbols (`==`, `!=`, etc.). Also, if you are doing numeric tests, keep in mind that `$x >= $y` is not the same as `$x => $y`. Be sure to use the correct operator!

Here is an example of a numeric test. If `$varname` is greater than 23, the code inside the curly braces is executed:

```
if ($varname > 23) {
    # do stuff here if the condition is true
}
```

If you need to have more than one condition, you can add `elsif` and `else` blocks:

```
if ($varname eq "somestring") {
    # do stuff here if the condition is true
}
elsif ($varname eq "someotherstring") {
    # do other stuff
}
else {
    # do this if none of the other conditions are met
}
```

The line breaks are not required; this example is just as valid:

```
if ($varname > 23) {
    print "$varname is greater than 23";
} elsif ($varname == 23) {
    print "$varname is 23";
} else { print "$varname is less than 23"; }
```

You can join conditions together by using logical operators:

Logical Operators

Operator	Example	Explanation
<code>&&</code>	<code>condition1 && condition2</code>	True if condition1 and condition2 are both true
<code> </code>	<code>condition1 condition2</code>	True if either condition1 or condition2 is true
<code>and</code>	<code>condition1 and condition2</code>	Same as <code>&&</code> but lower precedence
<code>or</code>	<code>condition1 or condition2</code>	Same as <code> </code> but lower precedence

Logical operators are evaluated from left to right. *Precedence* indicates which operator is evaluated first, in the event that more than one operator appears on one line. In a case like this:

```
condition1 || condition2 && condition3
```

`condition2 && condition3` is evaluated first, then the result of that evaluation is used in the `||` evaluation.

`and` and `or` work the same way as `&&` and `||`, although they have lower precedence than their symbolic counterparts.

Unless

`unless` is similar to `if`. Let's say you wanted to execute code only if a certain condition were false. You could do something like this:

```
if ($varname != 23) {
    # code to execute if $varname is not 23
}
```

The same test can be done using `unless`:

```
unless ($varname == 23) {
    # code to execute if $varname is not 23
}
```

There is no “elseunless”, but you can use an `else` clause:


```
unless ($varname == 23) {
    # code to execute if $varname is not 23
} else {
    # code to execute if $varname IS 23
}
```

Validating Form Data

You should always *validate* data submitted on a form; that is, check to see that the form fields aren't blank, and that the data submitted is in the format you expected. This is typically done with if/elsif blocks.

Here are some examples. This condition checks to see if the “name” field isn't blank:

```
if (param('name') eq "") {
    &dienice("Please fill out the field for your name.");
}
```

You can also test multiple fields at the same time:

```
if (param('name') eq "" or param('email') eq "") {
    &dienice("Please fill out the fields for your name
    and email address.");
}
```

The above code will return an error if either the name or email fields are left blank.

`param('fieldname')` always returns one of the following:

undef – or <i>undefined</i>	fieldname is not defined in the form itself, or it's a checkbox/radio button field that wasn't checked.
the empty string	fieldname exists in the form but the user didn't type anything into that field (for text fields)
one or more values	whatever the user typed into the field(s)

If your form has more than one field containing the same `fieldname`, then the values are stored sequentially in an array, accessed by `param('fieldname')`.

You should always validate all form data – even fields that are submitted as hidden fields

in your form. Don't assume that *your* form is always the one calling your program. *Any* external site can send data to your CGI. Never trust form input data.

Looping

Loops allow you to repeat code for as long as a condition is met. Perl has several loop control structures: `foreach`, `for`, `while` and `until`.

Foreach Loops

`foreach` iterates through a list of values:

```
foreach my $i (@arrayname) {  
    # code here  
}
```

This loops through each element of `@arrayname`, setting `$i` to the current array element for each pass through the loop. You may omit the loop variable `$i`:

```
foreach (@arrayname) {  
    # $_ is the current array element  
}
```

This sets the special Perl variable `$_` to each array element. `$_` does not need to be declared (it's part of the Perl language) and its scope localized to the loop itself.

For Loops

Perl also supports C-style `for` loops:

```
for ($i = 1; $i < 23; $i++) {  
    # code here  
}
```

The `for` statement uses a 3-part conditional: the loop initializer; the loop condition (how long to run the loop); and the loop re-initializer (what to do at the end of each iteration of the loop). In the above example, the loop initializes with `$i` being set to 1. The loop will run for as long as `$i` is less than 23, and at the end of each iteration `$i` is incremented by 1 using the auto-increment operator (`++`).

The conditional expressions are optional. You can do infinite loops by omitting all three conditions:

```
for (;;) {  
    # code here  
}
```

You can also write infinite loops with `while`.

While Loops

A while loop executes as long as particular condition is true:

```
while (condition) {  
    # code to run as long as condition is true  
}
```

Until Loops

`until` is the reverse of `while`. It executes as long as a particular condition is NOT true:

```
until (condition) {  
    # code to run as long as condition is not true  
}
```

Infinite Loops

An infinite loop is usually written like so:

```
while (1) {  
    # code here  
}
```

Obviously unless you want your program to run forever, you'll need some way to break out of these infinite loops. We'll look at breaking next.

Breaking from Loops

There are several ways to break from a loop. To stop the current loop iteration (and move on to the next one), use the `next` command:

```
foreach my $i (1..20) {  
    if ($i == 13) {  
        next;  
    }  
    print "$i\n";  
}
```

```
}
```

This example prints the numbers from 1 to 20, except for the number 13. When it reaches 13, it skips to the next iteration of the loop.

To break out of a loop entirely, use the `last` command:

```
foreach my $i (1..20) {  
    if ($i == 13) {  
        last;  
    }  
    print "$i\n";  
}
```

This example prints the numbers from 1 to 12, then terminates the loop when it reaches 13.

`next` and `last` only effect the innermost loop structure, so if you have something like this:

```
foreach my $i (@list1) {  
    foreach my $j (@list2) {  
        if ($i == 5 && $j == 23) {  
            last;  
        }  
    }  
    # this is where that last sends you  
}
```

The `last` command only terminates the innermost loop. If you want to break out of the outer loop, you need to use loop labels:

```
OUTER: foreach my $i (@list1) {  
    INNER: foreach my $j (@list2) {  
        if ($i == 5 && $j == 23) {  
            last OUTER;  
        }  
    }  
}  
# this is where that last sends you
```

The loop label is a string that appears before the loop command (`foreach`, `for`, or `while`). In this example we used `OUTER` as the label for the outer `foreach` loop and `INNER` for the inner loop label.

Now that you've seen the various types of Perl control structures, let's look at how to apply them to handling advanced form data.

Handling Checkboxes

Checkboxes allow the viewer to select one or more options on a form. If you assign each checkbox field a different name, you can print them the same way you'd print any form field using `param('fieldname')`.

Here is the HTML code for a set of checkboxes:

```
<b>Pick a Color:</b><br>
<form action="colors.cgi" method="POST">
<input type="checkbox" name="red" value=1> Red<br>
<input type="checkbox" name="green" value=1> Green<br>
<input type="checkbox" name="blue" value=1> Blue<br>
<input type="checkbox" name="gold" value=1> Gold<br>
<input type="submit">
</form>
```

⇒ Working example: <http://www.cgi101.com/book/ch5/colors.html>

This example lets the visitor pick as many options as they want – or none, if they prefer. Since this example uses a different field name for each checkbox, you can test it using `param`:

```
my @colors = ("red", "green", "blue", "gold");
foreach my $color (@colors) {
    if (param($color)) {
        print "You picked $color.\n";
    }
}
```

☞ Source code: <http://www.cgi101.com/book/ch5/colors-cgi.html>

Since we set the value of each checkbox to 1 (a true value), we didn't need to actually see if `param($color)` was equal to anything – if the box is checked, its true. If it's not checked, then `param($color)` is undefined and therefore not true.

The other way you could code this form is to set each checkbox name to the *same* name, and use a different value for each checkbox:

```
<b>Pick a Color:</b><br>

<form action="colors.cgi" method="POST">
<input type="checkbox" name="color" value="red"> Red<br>
<input type="checkbox" name="color" value="green"> Green<br>
<input type="checkbox" name="color" value="blue"> Blue<br>
<input type="checkbox" name="color" value="gold"> Gold<br>
<input type="submit">
</form>
```

⇒ Working example: <http://www.cgi101.com/book/ch5/colors2.html>

`param('color')` returns a list of the selected checkboxes, which you can then store in an array. Here is how you'd use it in your CGI program:

```
my @colors = param('color');
foreach my $color (@colors) {
    print "You picked $color.<br>\n";
}
```

☞ Source code: <http://www.cgi101.com/book/ch5/colors2-cgi.html>

Handling Radio Buttons

Radio buttons are similar to checkboxes in that you can have several buttons, but the difference is that the viewer can only pick one choice. As with our last checkbox example, the group of related radio buttons must all have the same name, and different values:

```
<b>Pick a Color:</b><br>

<form action="colors.cgi" method="POST">
<input type="radio" name="color" value="red"> Red<br>
<input type="radio" name="color" value="green"> Green<br>
<input type="radio" name="color" value="blue"> Blue<br>
<input type="radio" name="color" value="gold"> Gold<br>
<input type="submit">
</form>
```

⇒ Working example: <http://www.cgi101.com/book/ch5/colors3.html>

Since the viewer can only choose one item from a set of radio buttons, `param('color')` will be the color that was picked:

```
my $color = param('color');
print "You picked $color.<br>\n";
```

☞ Source code: <http://www.cgi101.com/book/ch5/colors3-cgi.html>

It's usually best to set the values of radio buttons to something meaningful; this allows you to print out the button name and its value, without having to store another list inside your CGI program. But if your buttons have lengthy values, or values unsuitable for storing in the value field, you can set each value to an abbreviation, then define a hash in your CGI program where the hash keys correspond to the abbreviations. The hash values can then contain longer data.

Let's try it. Create a new HTML form called colors4.html:

Program 5-1: colors4.html	Favorite Colors HTML Form
----------------------------------	----------------------------------

```
<html><head><title>Pick a Color</title></head>
<body>
<b>Pick a Color:</b><br>

<form action="colors4.cgi" method="POST">
<input type="radio" name="color" value="red"> Red<br>
<input type="radio" name="color" value="green"> Green<br>
<input type="radio" name="color" value="blue"> Blue<br>
<input type="radio" name="color" value="gold"> Gold<br>
<input type="submit">
</form>
</body>
</html>
```

⇒ Working example: <http://www.cgi101.com/book/ch5/colors4.html>

Next create colors4.cgi. This example not only prints out the color you picked, but also sets the page background to that color. The %colors hash stores the various RGB hex values for each color. The hex value for the selected color is then passed to CGI.pm's start_html function as the bgcolor (background color) parameter.

Program 5-2: colors4.cgi	Favorite Colors Program
---------------------------------	--------------------------------

```
#!/usr/bin/perl -wT
use strict;
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
```

```
my %colors = ( red    => "#ff0000",
               green  => "#00ff00",
               blue   => "#0000ff",
               gold   => "#cccc00");

print header;
my $color = param('color');

# do some validation - be sure they picked a valid color
if (exists $colors{$color}) {
    print start_html(-title=>"Results", -bgcolor=>$color);
    print "You picked $color.<br>\n";
} else {
    print start_html(-title=>"Results");
    print "You didn't pick a color! (You picked '$color')";
}
print end_html;
```

☞ Source code: <http://www.cgi101.com/book/ch5/colors4-cgi.html>

Handling SELECT Fields

SELECT fields are handled almost the same way as radio buttons. A SELECT field is a pull-down menu with one or more choices. Unless you specify a multiple select (see below), the viewer can only choose one option. Here is the HTML for creating a SELECT field:

```
<select name="color">
<option value="red"> Red
<option value="green"> Green
<option value="blue"> Blue
<option value="gold"> Gold
</select>
```

As with radio buttons, you access the selection in your CGI program using `param('color')`:

```
my $color = param('color');
print "You picked $color.<br>\n";
```

Multiple-choice SELECTs

Multiple SELECTs allow the viewer to choose more than one option from the list, usually

by option-clicking or control-clicking on the options they want. Here is the HTML for a multiple SELECT:

```
<select name="color" multiple size=3>
<option value="red"> Red
<option value="green"> Green
<option value="blue"> Blue
<option value="gold"> Gold
</select>
```

In your CGI program, `param('color')` returns a list of the selected values, just as it did when we had multiple checkboxes of the same name:

```
my @colors = param('color');
foreach my $color (@colors) {
    print "You picked $color.<br>\n";
}
```

So now you've seen every type of form element (except for file-uploads, which we'll look at in Chapter 14), and in every case you've seen that CGI.pm's `param` function returns the value (or values) from each form field. The value returned by `param` is always a list, but for text, textarea, password, radio, and single select fields you can use it in a scalar context. For checkboxes and multiple select fields, you use it in an array context.

In the next chapter we'll learn how to read and write data files, so you'll be able to save and analyze the data collected by your forms.

Resources

Visit <http://www.cgi101.com/book/ch5/> for source code and links from this chapter.



Reading and Writing Data Files

As you start to program more advanced CGI applications, you'll want to store data so you can use it later. Maybe you have a guestbook program and want to keep a log of the names and email addresses of visitors, or a page counter that must update a counter file, or a program that scans a flat-file database and draws info from it to generate a page. You can do this by reading and writing data files (often called *file I/O*).

File Permissions

Most web servers run with very limited permissions; this protects the server (and the system it's running on) from malicious attacks by users or web visitors. On Unix systems, the web process runs under its own userid, typically the "web" or "nobody" user. Unfortunately this means the server doesn't have permission to create files in *your* directory. In order to write to a data file, you must usually make the file (or the directory where the file will be created) world-writable – or at least writable by the web process userid. In Unix a file can be made world-writable using the **chmod** command:

```
chmod 666 myfile.dat
```

To set a directory world-writable, you'd do:

```
chmod 777 directoryname
```

See Appendix A for a chart of the various **chmod** permissions.

Unfortunately, if the file is world-writable, it can be written to (or even deleted) by other users on the system. You should be very cautious about creating world-writable files in your web space, and you should *never* create a world-writable directory there. (An attacker could use this to install their own CGI programs there.) If you must have a world-writable directory, either use /tmp (on Unix), or a directory outside of your web

space. For example if your web pages are in `/home/you/public_html`, set up your writable files and directories in `/home/you`.

A much better solution is to configure the server to run your programs with your `userid`. Some examples of this are `CGIwrap` (platform independent) and `suEXEC` (for Apache/Unix). Both of these force CGI programs on the web server to run under the program *owner's* `userid` and permissions. Obviously if your CGI program is running with your `userid`, it will be able to create, read and write files in your directory without needing the files to be world-writable.

The Apache web server also allows the webmaster to define what user and group the server runs under. If you have your own domain, ask your webmaster to set up your domain to run under your own `userid` and group permissions.

Permissions are less of a problem if you only want to *read* a file. If you set the file permissions so that it is group- and world-readable, your CGI programs can then safely read from that file. Use caution, though; if your program can read the file, so can the webserver, and if the file is in your webspace, someone can type the direct URL and view the contents of the file. Be sure not to put sensitive data in a publicly readable file.

Opening Files

Reading and writing files is done by opening a file and associating it with a *filehandle*. This is done with the statement:

```
open(filehandle, filename);
```

The filename may be prefixed with a `>`, which means to overwrite anything that's in the file now, or with a `>>`, which means to append to the bottom of the existing file. If both `>` and `>>` are omitted, the file is opened for reading only. Here are some examples:

```
open(INF, "out.txt");           # opens mydata.txt for reading
open(OUTF, ">out.txt");         # opens out.txt for overwriting
open(OUTF, ">>out.txt");        # opens out.txt for appending
open(FH, "+<out.txt");         # opens existing file out.txt for
                                #   reading AND writing
```

The filehandles in these cases are `INF`, `OUTF` and `FH`. You can use just about any name for the filehandle.

Also, a warning: your web server might do strange things with the path your programs run under, so it's possible you'll have to use the full path to the file (such as

/home/you/public_html/somedata.txt), rather than just the filename. This is generally not the case with the Apache web server, but some other servers behave differently. Try opening files with just the filename first (provided the file is in the same directory as your CGI program), and if it doesn't work, then use the full path.

One problem with the above code is that it doesn't check the return value of `open` to ensure the file was really opened. `open` returns nonzero upon success, or `undef` (which is a false value) otherwise. The safe way to open a file is as follows:

```
open(OUTF,">outdata.txt") or &dienice("Can't open
outdata.txt for writing: $!");
```

This uses the “dienice” subroutine we wrote in Chapter 4 to display an error message and exit if the file can't be opened. You should do this for all file opens, because if you don't, your CGI program will continue running even if the file isn't open, and you could end up losing data. It can be quite frustrating to realize you've had a survey running for several weeks while no data was being saved to the output file.

The `$!` in the above example is a special Perl variable that stores the error code returned by the failed `open` statement. Printing it may help you figure out why the open failed.

Guestbook Form with File Write

Let's try this by modifying the guestbook program you wrote in Chapter 4. The program already sends you e-mail with the information; we're going to have it write its data to a file as well.

First you'll need to create the output file and make it writable, because your CGI program probably can't create new files in your directory. If you're using Unix, log into the Unix shell, **cd** to the directory where your guestbook program is located, and type the following:

```
touch guestbook.txt
chmod 622 guestbook.txt
```

The Unix **touch** command, in this case, creates a new, empty file called “guestbook.txt”. (If the file already exists, **touch** simply updates the last-modified timestamp of the file.) The **chmod 622** command makes the file read/write for you (the owner), and write-only for everyone else.

If you don't have Unix shell access (or you aren't using a Unix system), you should create or upload an empty file called `guestbook.txt` in the directory where your

guestbook.cgi program is located, then adjust the file permissions on it using your FTP program.

Now you'll need to modify guestbook.cgi to write to the file:

Program 6-1: guestbook.cgi**Guestbook Program With File Write**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;

print header;
print start_html("Results");

# first print the mail message...

$ENV{PATH} = "/usr/sbin";
open (MAIL, "|/usr/sbin/sendmail -oi -t -odq") or
    &dienice("Can't fork for sendmail: $!\n");
print MAIL "To: recipient\@cgil01.com\n";
print MAIL "From: nobody\@cgil01.com\n";
print MAIL "Subject: Form Data\n\n";
foreach my $p (param()) {
    print MAIL "$p = ", param($p), "\n";
}
close(MAIL);

# now write (append) to the file

open(OUT, ">>guestbook.txt") or &dienice("Couldn't open
output file: $!");
foreach my $p (param()) {
    print OUT param($p), "|";
}
print OUT "\n";
close(OUT);

print <<EndHTML;
<h2>Thank You</h2>
<p>Thank you for writing!</p>
<p>Return to our <a href="index.html">home page</a>.</p>
EndHTML

print end_html;
```

```

sub dienice {
    my($errmsg) = @_;
    print "<h2>Error</h2>\n";
    print "<p>$errmsg</p>\n";
    print end_html;
    exit;
}

```

☞ Source code: <http://www.cgi101.com/book/ch6/guestbook-cgi.html>

⇒ Working example: <http://www.cgi101.com/book/ch6/guestbook.html>

Now go back to your browser and fill out the guestbook form again. If your CGI program runs without any errors, you should see data added to the `guestbook.txt` file. The resulting file will show the submitted form data in pipe-separated form:

```
Someone|someone@wherever.com|comments here
```

Ideally you'll have one line of data (or *record*) for each form that is filled out. This is what's called a *flat-file database*.

Unfortunately if the visitor enters multiple lines in the comments field, you'll end up with multiple lines in the data file. To remove the newlines, you should substitute newline characters (`\n`) as well as hard returns (`\r`). Perl has powerful pattern matching and replacement capabilities; it can match the most complex patterns in a string using *regular expressions* (see Chapter 13). The basic syntax for substitution is:

```
$mystring =~ s/pattern/replacement/;
```

This command substitutes “pattern” for “replacement” in the scalar variable `$mystring`. Notice the operator is a `~` (an equals sign followed by a tilde); this is Perl's *binding operator* and indicates a regular expression pattern match/substitution/replacement is about to follow.

Here is how to replace the end-of-line characters in your guestbook program:

```

foreach my $p (param()) {
    my $value = param($p);
    $value =~ s/\n/ /g;      # replace newlines with spaces
    $value =~ s/\r//g;       # remove hard returns
    print OUT "$p = $value,";
}

```

Go ahead and change your program, then test it again in your browser. View the

guestbook.txt file in your browser or in a text editor and observe the results.

File Locking

CGI processes on a Unix web server can run simultaneously, and if two programs try to open and write the same file at the same time, the file may be erased, and you'll lose all of your data. To prevent this, you need to lock the files you are writing to. There are two types of file locks:

- A *shared lock* allows more than one program (or other process) to access the file at the same time. A program should use a shared lock when reading from a file.
- An *exclusive lock* allows only one program or process to access the file while the lock is held. A program should use an exclusive lock when writing to a file.

File locking is accomplished in Perl using the Fcntl module (which is part of the standard library), and the `flock` function. The `use` statement is like CGI.pm's:

```
use Fcntl qw(:flock);
```

The Fcntl module provides *symbolic values* (like abbreviations) representing the correct lock numbers for the `flock` function, but you must specify `:flock` in the `use` statement in order for Fcntl to export those values. The values are as follows:

LOCK_SH	shared lock
LOCK_EX	exclusive lock
LOCK_NB	non-blocking lock
LOCK_UN	unlock

These abbreviations can then be passed to `flock`. The `flock` function takes two arguments: the filehandle and the lock type, which is typically a number. The number may vary depending on what operating system you are using, so it's best to use the symbolic values provided by Fcntl. A file is locked *after* you open it (because the filehandle doesn't exist before you open the file):

```
open(FH, "filename") or &dienice("Can't open file: $!");  
flock(FH, LOCK_SH);
```

The lock will be released automatically when you close the file or when the program finishes.

Keep in mind that file locking is only effective if *all* of the programs that read and write

to that file also use `flock`. Programs that don't will ignore the locks held by other processes.

Since `flock` may force your CGI program to wait for another process to finish writing to a file, you should also reset the file pointer, using the `seek` function:

```
seek(filehandle, offset, whence);
```

`offset` is the number of bytes to move the pointer, relative to `whence`, which is one of the following:

0	beginning of file
1	current file position
2	end of file

So `seek(OUTF, 0, 2)` repositions the pointer to the end of the file. If you were reading the file instead of writing to it, you'd want to do `seek(OUTF, 0, 0)` to reset the pointer to the beginning of the file.

The `Fcntl` module also provides symbolic values for the seek pointers:

<code>SEEK_SET</code>	beginning of file
<code>SEEK_CUR</code>	current file position
<code>SEEK_END</code>	end of file

To use these, add `:seek` to the `use Fcntl` statement:

```
use Fcntl qw(:flock :seek);
```

Now you can use `seek(OUTF, 0, SEEK_END)` to reset the file pointer to the end of the file, or `seek(OUTF, 0, SEEK_SET)` to reset it to the beginning of the file.

Closing Files

When you're finished writing to a file, it's best to close the file, like so:

```
close(filehandle);
```

Files are automatically closed when your program ends. File locks are released when the file is closed, so it is not necessary to actually unlock the file before closing it. (In fact, releasing the lock before the file is closed can be dangerous and cause you to lose data.)

Reading Files

There are two ways you can handle reading data from a file: you can either read one line at a time, or read the entire file into an array. Here's an example:

```
open(FH,"guestbook.txt") or &dienice("Can't open
guestbook.txt: $!");

my $a = <FH>;      # reads one line from the file into
                   # the scalar $a
my @b = <FH>;      # reads the ENTIRE FILE into array @b

close(FH);         # closes the file
```

If you were to use this code in your program, you'd end up with the first line of `guestbook.txt` being stored in `$a`, and the remainder of the file in array `@b` (with each element of `@b` containing one line of data from the file). The actual read occurs with `<filehandle>`; the amount of data read depends on the type of variable you save it into.

The following section of code shows how to read the entire file into an array, then loop through each element of the array to print out each line:

```
open(FH,"guestbook.txt") or &dienice("Can't open
guestbook.txt: $!");
my @ary = <FH>;
close(FH);

foreach my $line (@ary) {
    print $line;
}
```

This code minimizes the amount of time the file is actually open. The drawback is it causes your CGI program to consume as much memory as the size of the file. Obviously for very large files that's not a good idea; if your program consumes more memory than the machine has available, it could crash the whole machine (or at the very least make things extremely slow). To process data from a very large file, it's better to use a `while` loop to read one line at a time:

```
open(FH,"guestbook.txt") or &dienice("Can't open
guestbook.txt: $!");
while (my $line = <FH>) {
    print $line;
}
close(FH);
```

Poll Program

Let's try another example: a web poll. You've probably seen them on various news sites. A basic poll consists of one question and several potential answers (as radio buttons); you pick one of the answers, vote, then see the poll results on the next page.

Start by creating the poll HTML form. Use whatever question and answer set you wish.

Program 6-2: poll.html**Poll HTML Form**

```
<form action="poll.cgi" method="POST">
Which was your favorite <i>Lord of the Rings</i> film?<br>
<input type="radio" name="pick" value="fotr">The Fellowship
of the Ring<br>
<input type="radio" name="pick" value="ttt">The Two
Towers<br>
<input type="radio" name="pick" value="rotk">Return of the
King<br>
<input type="radio" name="pick" value="none">I didn't watch
them<br>
<input type="submit" value="Vote">
</form>
<a href="results.cgi">View Results</a><br>
```

⇒ Working example: <http://www.cgi101.com/book/ch6/poll.html>

In this example we're using abbreviations for the radio button values. Our CGI program will translate the abbreviations appropriately.

Now the voting CGI program will write the result to a file. Rather than having this program analyze the results, we'll simply use a redirect to bounce the viewer to a third program (results.cgi). That way you won't need to write the results code twice.

Here is how the voting program (poll.cgi) should look:

Program 6-3: poll.cgi**Poll Program**

```
#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;
use Fcntl qw(:flock :seek);
```

```

my $outfile = "poll.out";

# only record the vote if they actually picked something
if (param('pick')) {
    open(OUT, ">>$outfile") or &dienice("Couldn't open
$outfile: $!");
    flock(OUT, LOCK_EX);      # set an exclusive lock
    seek(OUT, 0, SEEK_END);   # then seek the end of file
    print OUT param('pick'), "\n";
    close(OUT);
} else {
    # this is optional, but if they didn't vote, you might
    # want to tell them about it...
    &dienice("You didn't pick anything!");
}

# redirect to the results.cgi.
# (Change to your own URL...)
print redirect("http://cgi101.com/book/ch6/results.cgi");

sub dienice {
    my($msg) = @_ ;
    print header;
    print start_html("Error");
    print h2("Error");
    print $msg;
    print end_html;
    exit;
}

```

☞ Source code: <http://www.cgi101.com/book/ch6/poll-cgi.html>

Finally results.cgi reads the file where the votes are stored, totals the overall votes as well as the votes for each choice, and displays them in table format.

Program 6-4: results.cgi	Poll Results Program
---------------------------------	-----------------------------

```

#!/usr/bin/perl -wT
use CGI qw(:standard);
use CGI::Carp qw(warningsToBrowser fatalsToBrowser);
use strict;
use Fcntl qw(:flock :seek);

my $outfile = "poll.out";

```

```

print header;
print start_html("Results");

# open the file for reading
open(IN, "$outfile") or &dienice("Couldn't open $outfile:
$!");
# set a shared lock
flock(IN, LOCK_SH);
# then seek the beginning of the file
seek(IN, 0, SEEK_SET);

# declare the totals variables
my($total_votes, %results);
# initialize all of the counts to zero:
foreach my $i ("fotr", "ttt", "rotk", "none") {
    $results{$i} = 0;
}

# now read the file one line at a time:
while (my $rec = <IN>) {
    chomp($rec);
    $total_votes = $total_votes + 1;
    $results{$rec} = $results{$rec} + 1;
}
close(IN);

# now display a summary:
print <<End;
<b>Which was your favorite <i>Lord of the Rings</i> film?
</b><br>
<table border=0 width=50%>
<tr>
    <td>The Fellowship of the Ring</td>
    <td>$results{fotr} votes</td>
</tr>
<tr>
    <td>The Two Towers</td>
    <td>$results{ttt} votes</td>
</tr>
<tr>
    <td>Return of the King</td>
    <td>$results{rotk} votes</td>
</tr>
<tr>
    <td>didn't watch them</td>
    <td>$results{none} votes</td>
</tr>

```

```
</table>
<p>
$total_votes votes total
</p>
End

print end_html;

sub dienice {
    my($msg) = @_;
    print h2("Error");
    print $msg;
    print end_html;
    exit;
}
```

☞ Source code: <http://www.cgi101.com/book/ch6/results-cgi.html>

The results program only shows the total number of votes. You may also want to calculate the percentages and display a bar-graph for each vote relative to the overall total. We'll look at how to calculate percentages in the next chapter.

Resources

CGIwrap: <http://sourceforge.net/projects/cgiwrap>

Visit <http://www.cgi101.com/book/ch6/> for source code and links from this chapter.

Order *CGI Programming 101* Today!

The rest of the book is packed with more great CGI programming information and examples:

Chapter 7: Working With Strings: Comparing, replacing, reversing, and splitting strings; changing case

Chapter 8: Server-Side Includes: a complete list of Apache SSI directives; how to include files and execute CGI programs using SSIs

Chapter 9: Working with Numbers: Numeric functions and operators, random numbers, a random image program, a banner ad program

Chapter 10: Redirection: Counting clicks from banner ads; redirection based on referring page or visitor's country; multi-site redirector

Chapter 11: Multi-Script Forms: build an online catalog and order form system

Chapter 12: Searching and Sorting: Searching with `grep`; searching for multiple keywords; custom sorting subroutines

Chapter 13: Regular Expressions: Pattern matching and replacement

Chapter 14: Perl Modules: How to find and install modules (for both Unix and Windows), validating e-mail addresses with `Email::Valid`; uploading files from a form; a graphical counter program; e-mailing attachments with `MIME::Lite`

Chapter 15: Date and Time: Calculating and formatting dates and times; a countdown clock program

Chapter 16: Database Programming: Working with MySQL and the Perl DBI (database independent) module.

Chapter 17: HTTP Cookies: How to set, read, delete and track cookies; develop a cookie-based shopping cart

Chapter 18: Writing Your Own Modules: How to write a module for sharing common code.

Chapter 19: CGI Security: Protecting your site from hackers and spammers

Chapter 20: Password Protection: HTTP-authentication and cookie-based authentication; user registration program; password change programs; logout program

Appendix A: Unix Tutorial and Command Reference

A full-length PDF of *CGI Programming 101* is available from **e23**
(<http://e23.sjgames.com/item.html?id=cgi101>)

The printed book can be ordered from Amazon.com, or directly from the author at <http://www.cgi101.com/book/order.html>